Step-by-Step Evaluation Metrics

Model: meta-llama/Meta-Llama-3-8B | Dataset: Financial_PhraseBank | N: 16

Acc: 0.3223 | Loss: 1.0707 | Latency: 7.51ms | VRAM: 14.28GB

Model: meta-llama/Meta-Llama-3-8B | Dataset: Financial_PhraseBank | N: 32

Acc: 0.3333 | Loss: 1.0644 | Latency: 7.51ms | VRAM: 14.28GB

Model: meta-llama/Meta-Llama-3-8B | Dataset: Financial_PhraseBank | N: 64

Acc: 0.4812 | Loss: 0.9701 | Latency: 7.52ms | VRAM: 14.29GB

Model: meta-llama/Meta-Llama-3-8B | Dataset: Financial_PhraseBank | N: 128

Acc: 0.4327 | Loss: 1.0377 | Latency: 7.51ms | VRAM: 14.29GB

Model: meta-llama/Meta-Llama-3-8B | Dataset: Financial_PhraseBank | N: 256

Acc: 0.6336 | Loss: 0.8154 | Latency: 7.53ms | VRAM: 14.29GB

Model: meta-llama/Meta-Llama-3-8B | Dataset: SuperGLUE_RTE | N: 16

Acc: 0.4558 | Loss: 0.7407 | Latency: 7.53ms | VRAM: 14.29GB

Model: meta-llama/Meta-Llama-3-8B | Dataset: SuperGLUE_RTE | N: 32

Acc: 0.4779 | Loss: 0.6670 | Latency: 7.53ms | VRAM: 14.29GB

Model: meta-llama/Meta-Llama-3-8B | Dataset: SuperGLUE_RTE | N: 64

Acc: 0.4940 | Loss: 0.6874 | Latency: 7.52ms | VRAM: 14.29GB

Model: meta-llama/Meta-Llama-3-8B | Dataset: SuperGLUE_RTE | N: 128

Acc: 0.4960 | Loss: 0.6995 | Latency: 7.53ms | VRAM: 14.29GB

Model: meta-llama/Meta-Llama-3-8B | Dataset: SuperGLUE_RTE | N: 256

Acc: 0.5321 | Loss: 0.6732 | Latency: 7.54ms | VRAM: 14.29GB

Model: Qwen/Qwen3-8B | Dataset: Financial_PhraseBank | N: 16

Acc: 0.4658 | Loss: 0.2102 | Latency: 8.51ms | VRAM: 14.39GB

Model: Qwen/Qwen3-8B | Dataset: Financial_PhraseBank | N: 32

Acc: 0.5430 | Loss: 0.1474 | Latency: 8.51ms | VRAM: 14.39GB

Model: Qwen/Qwen3-8B | Dataset: Financial_PhraseBank | N: 64

Acc: 0.5342 | Loss: 0.3996 | Latency: 8.52ms | VRAM: 14.39GB

Model: Qwen/Qwen3-8B | Dataset: Financial_PhraseBank | N: 128

Acc: 0.5673 | Loss: 0.5977 | Latency: 8.53ms | VRAM: 14.39GB

Model: Qwen/Qwen3-8B | Dataset: Financial_PhraseBank | N: 256

Acc: 0.6049 | Loss: 0.0773 | Latency: 8.53ms | VRAM: 14.39GB

Model: Qwen/Qwen3-8B | Dataset: SuperGLUE_RTE | N: 16

Acc: 0.4558 | Loss: 0.0450 | Latency: 8.53ms | VRAM: 14.39GB

Model: Qwen/Qwen3-8B | Dataset: SuperGLUE_RTE | N: 32

Acc: 0.4859 | Loss: 0.0118 | Latency: 8.54ms | VRAM: 14.39GB

Model: Qwen/Qwen3-8B | Dataset: SuperGLUE_RTE | N: 64

Acc: 0.4759 | Loss: 0.2412 | Latency: 8.52ms | VRAM: 14.39GB

Model: Qwen/Qwen3-8B | Dataset: SuperGLUE_RTE | N: 128

Acc: 0.5301 | Loss: 0.0005 | Latency: 8.52ms | VRAM: 14.39GB

Model: Qwen/Qwen3-8B | Dataset: SuperGLUE_RTE | N: 256

Acc: 0.5361 | Loss: 0.5247 | Latency: 8.54ms | VRAM: 14.39GB

Tradeoff Analysis: Financial_PhraseBank - Accuracy

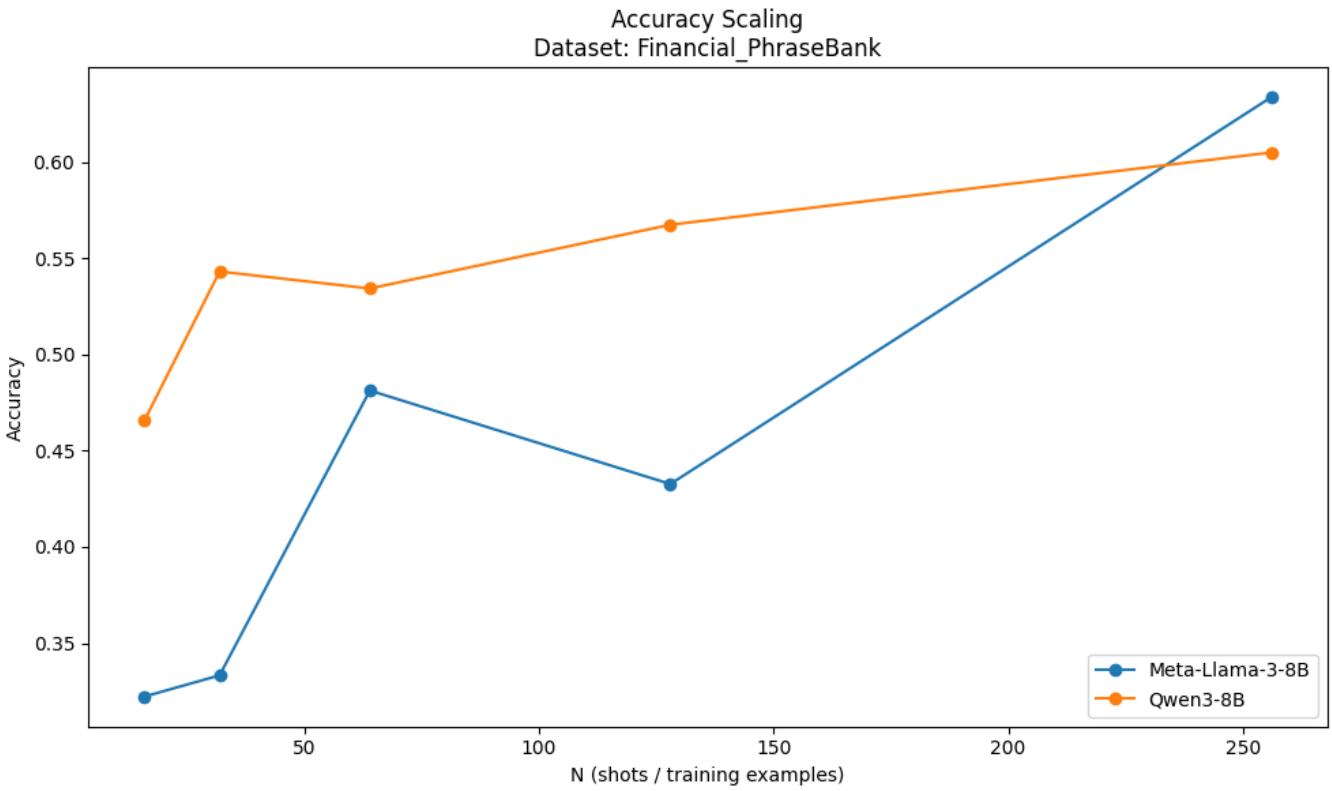

Tradeoff Analysis: Financial_PhraseBank - Final Training Loss

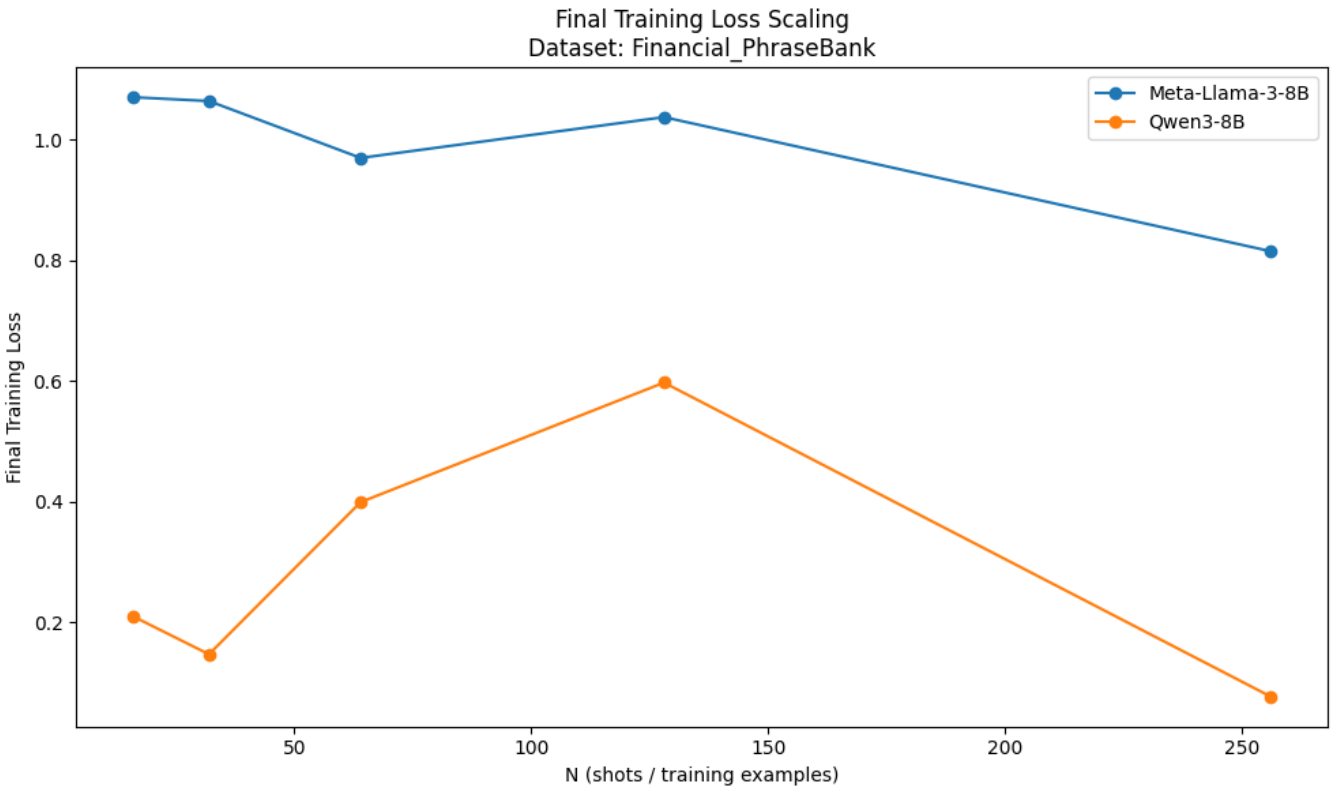

Tradeoff Analysis: Financial_PhraseBank - Latency (ms)

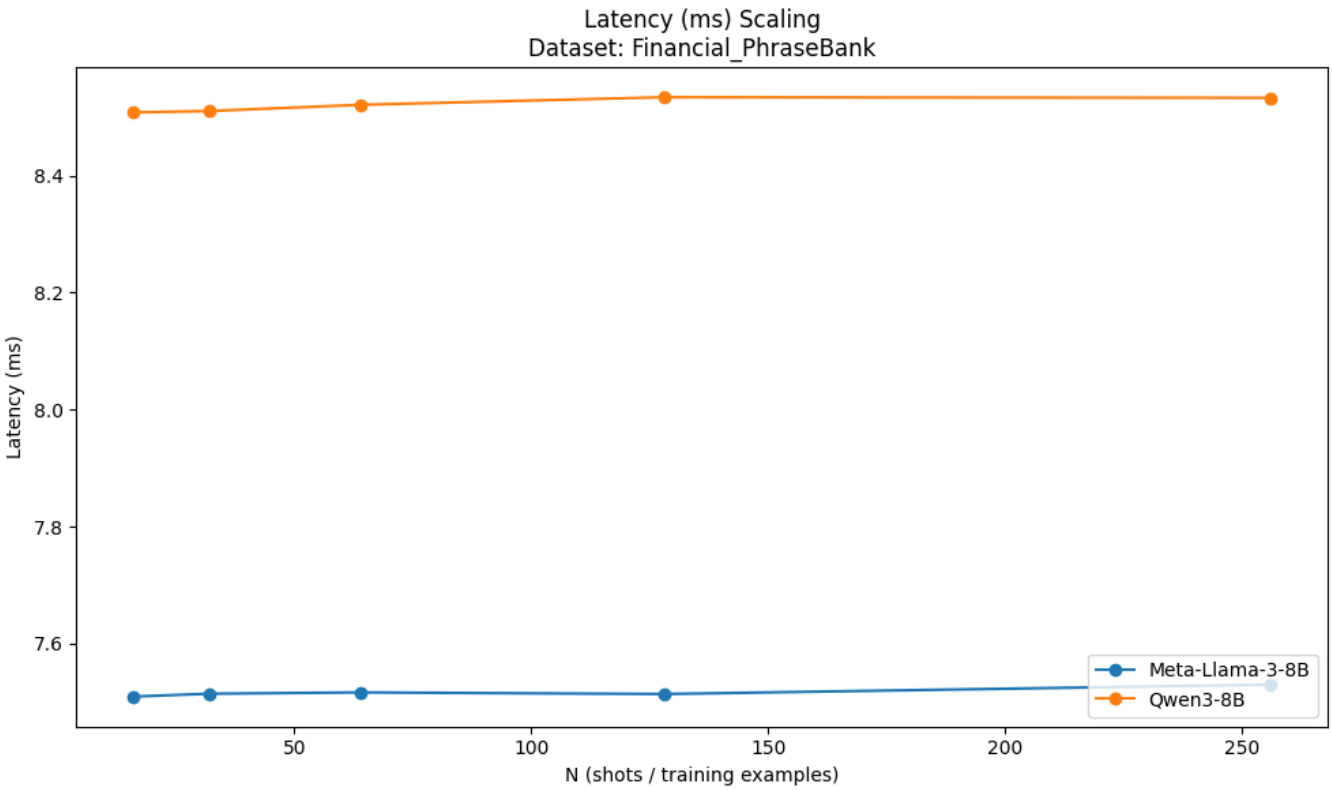

Tradeoff Analysis: Financial_PhraseBank - Throughput (samples/sec)

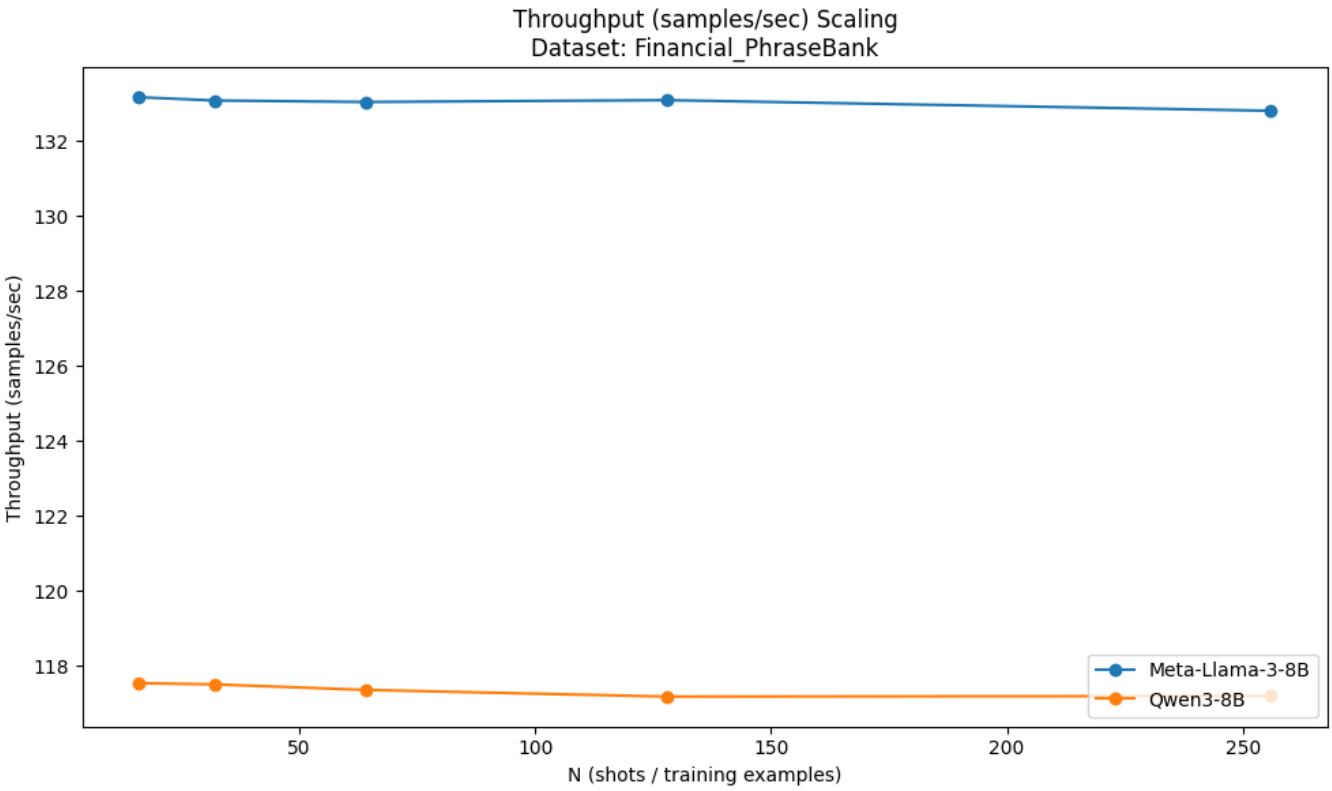

Tradeoff Analysis: Financial_PhraseBank - Peak VRAM (GB)

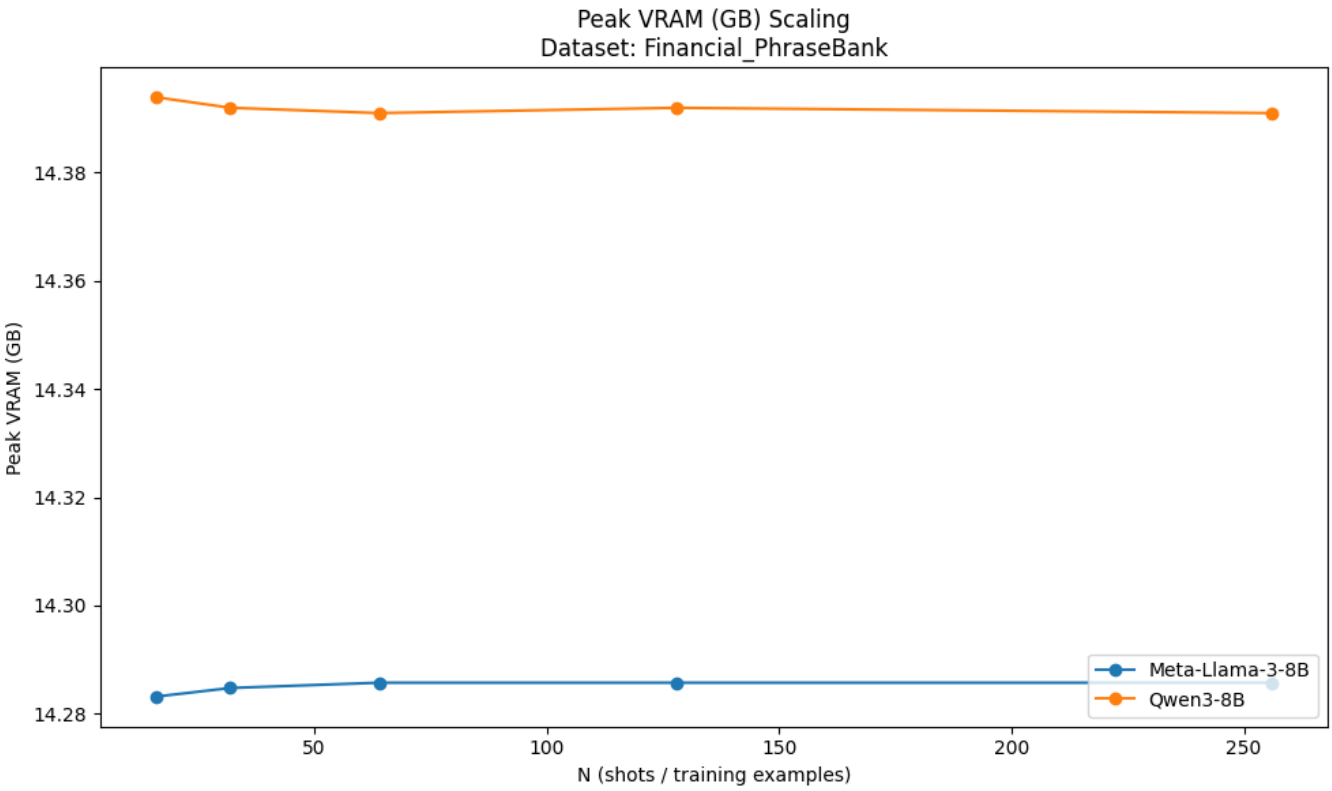

Tradeoff Analysis: SuperGLUE_RTE - Accuracy

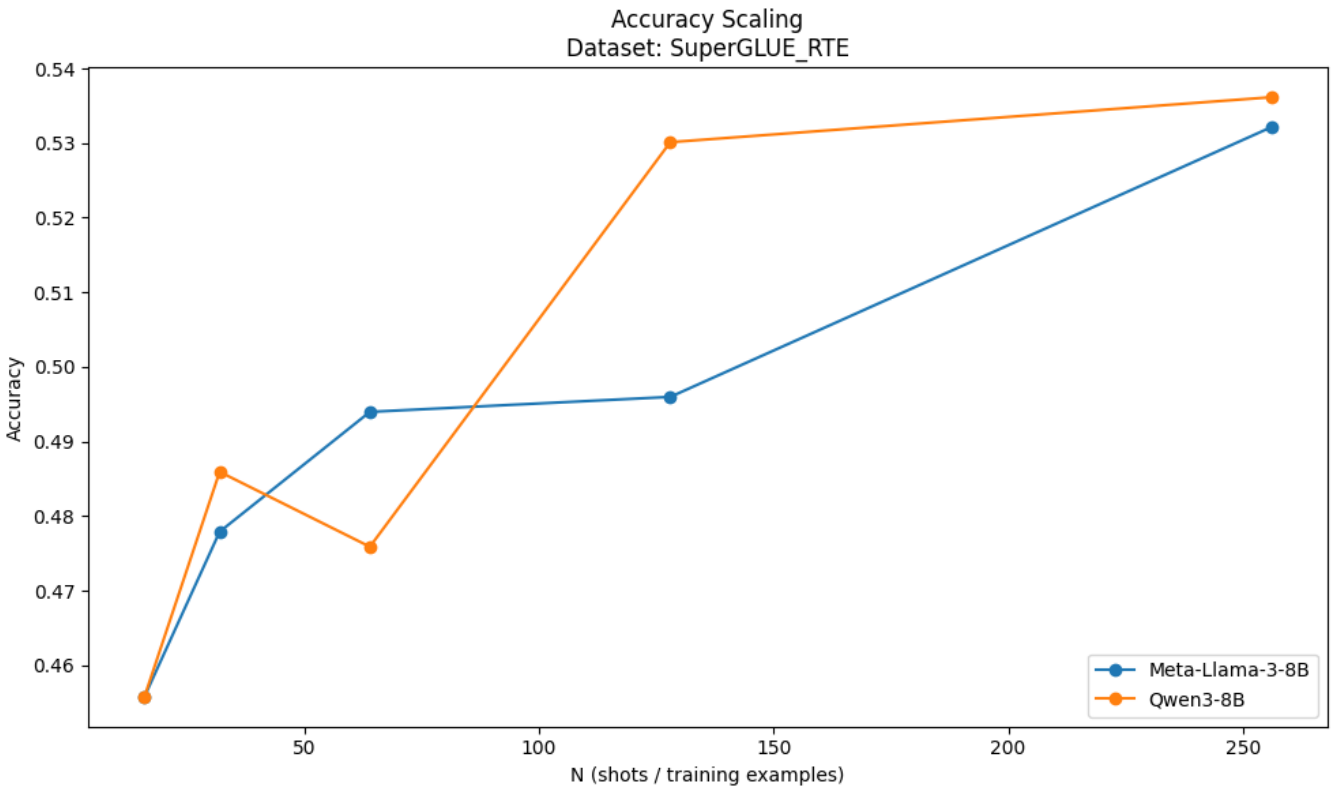

Tradeoff Analysis: SuperGLUE_RTE - Final Training Loss

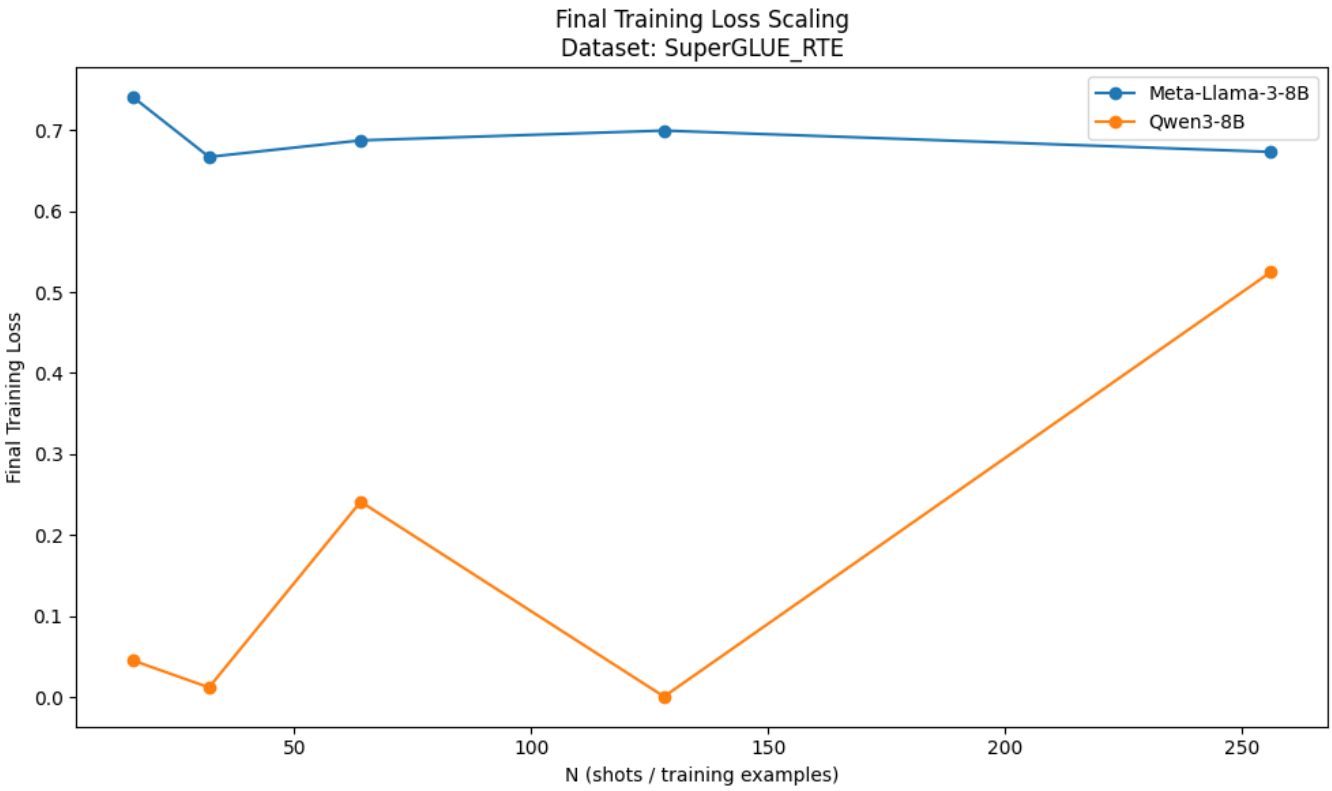

Tradeoff Analysis: SuperGLUE_RTE - Latency (ms)

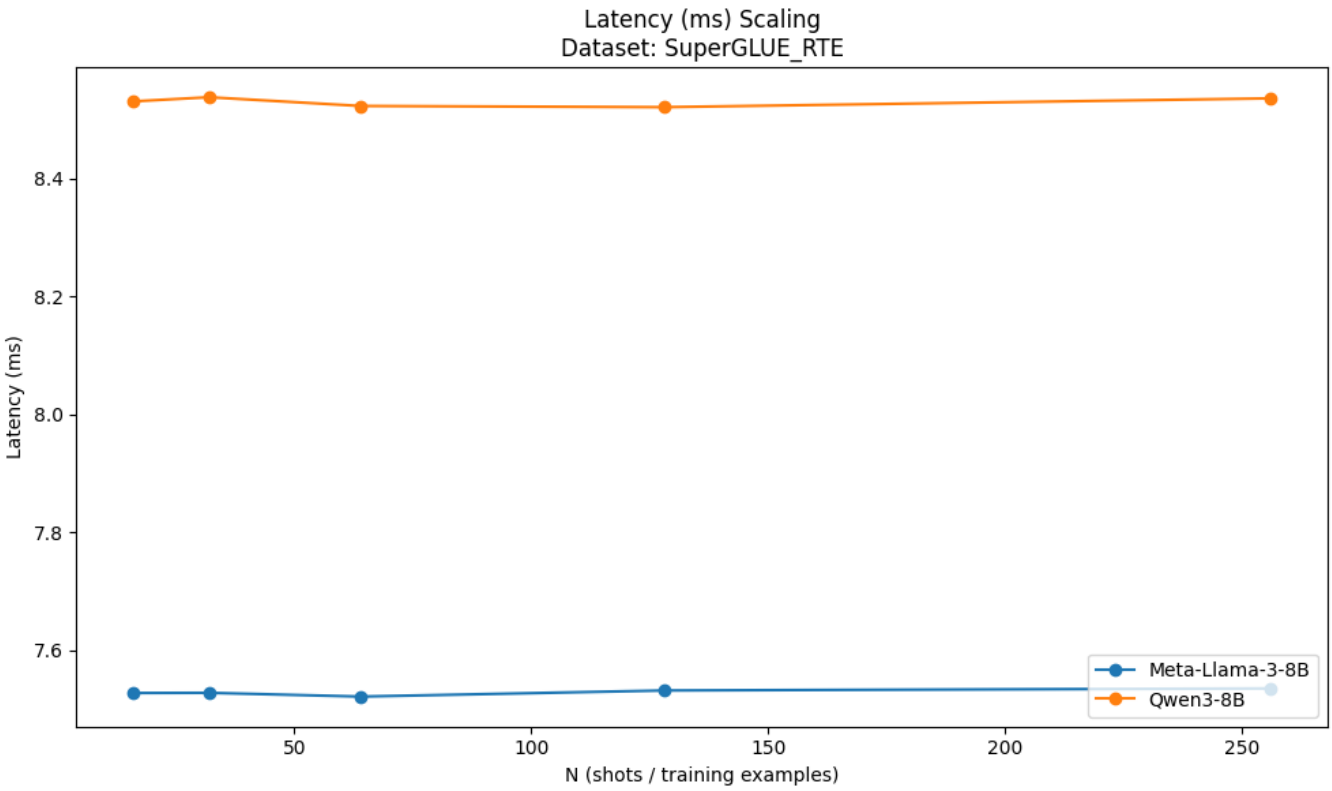

Tradeoff Analysis: SuperGLUE_RTE - Throughput (samples/sec)

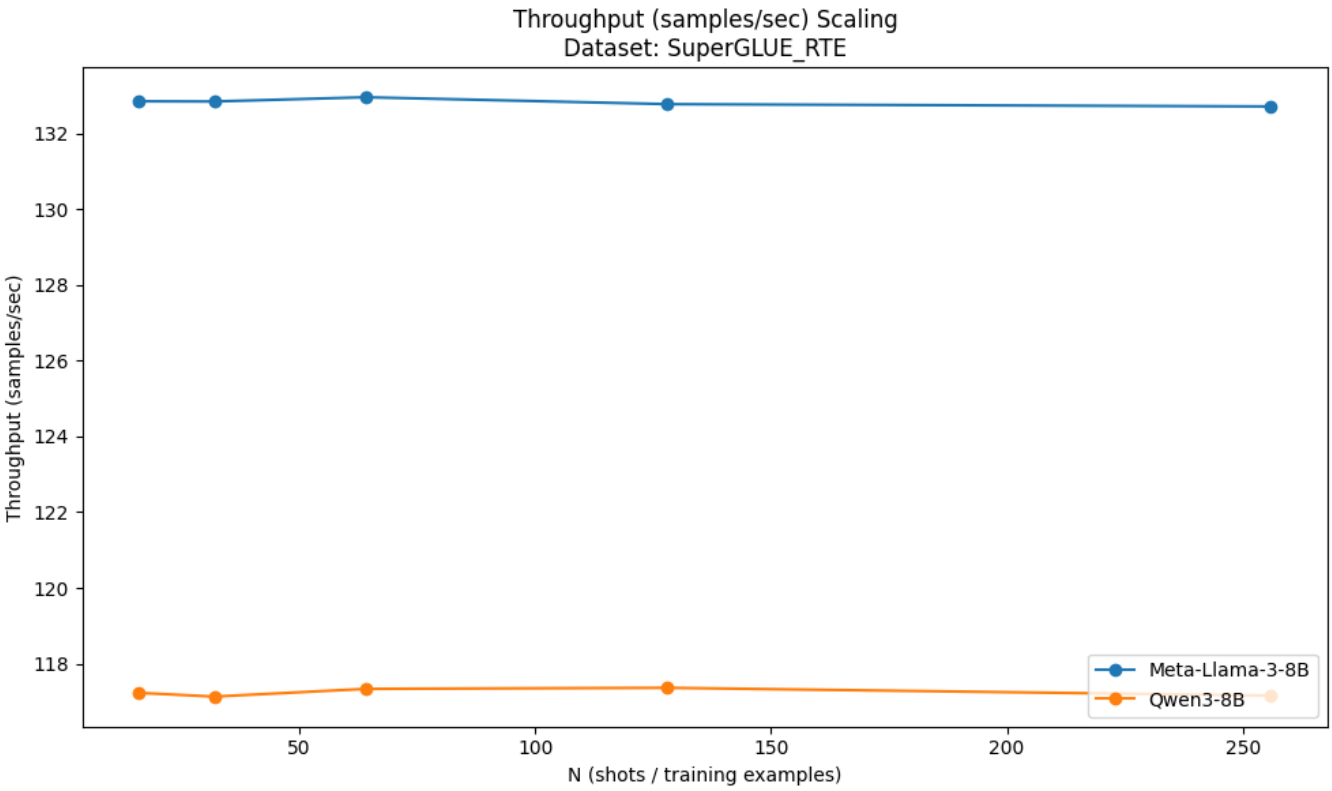

Tradeoff Analysis: SuperGLUE_RTE - Peak VRAM (GB)

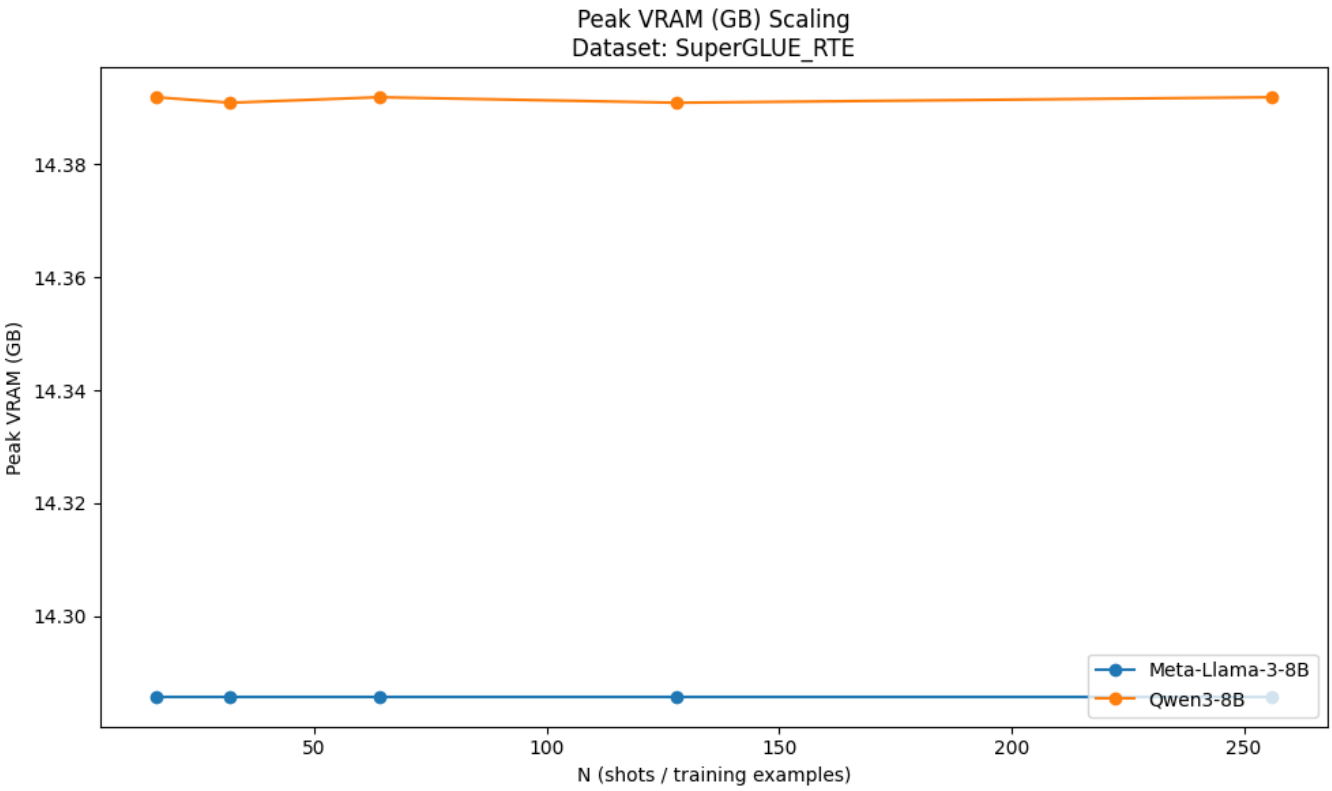

Training Loss Curves

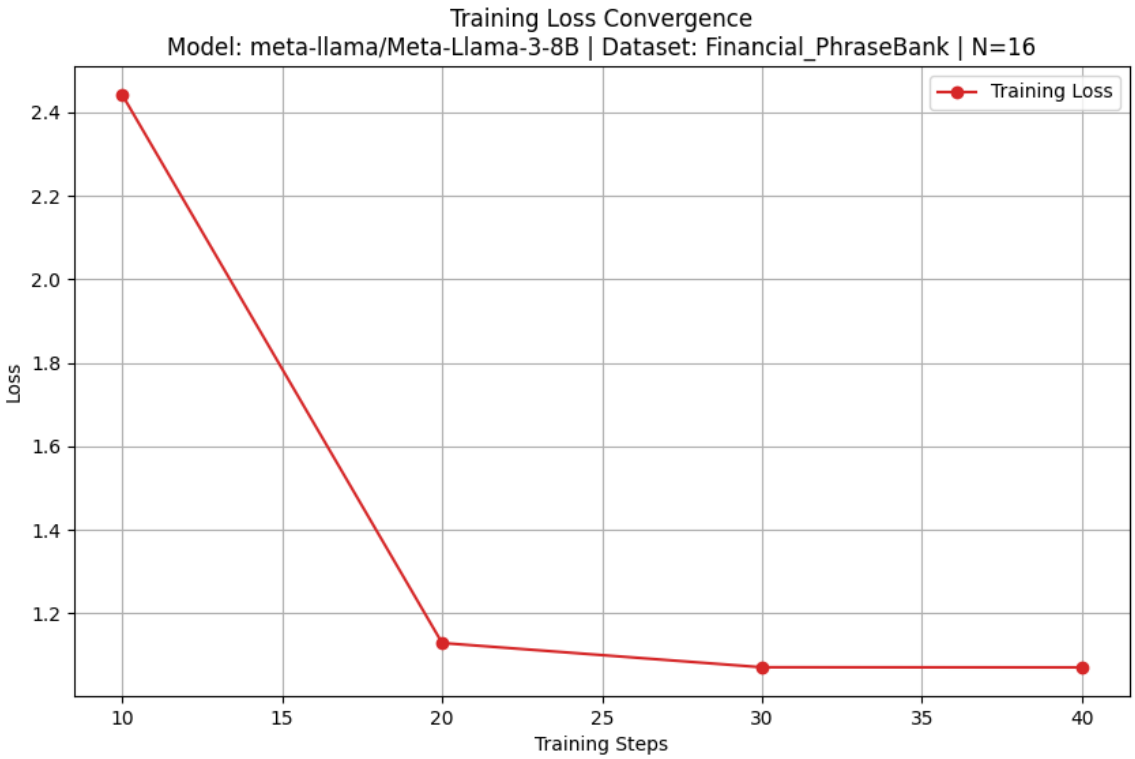

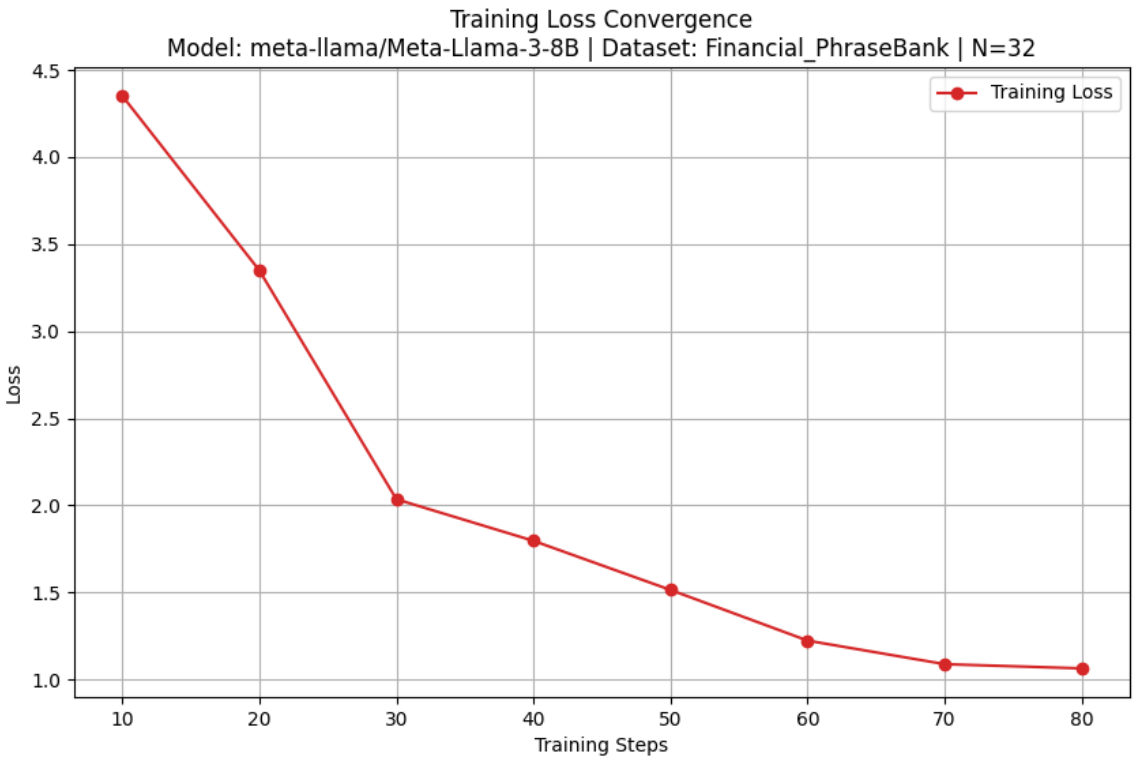

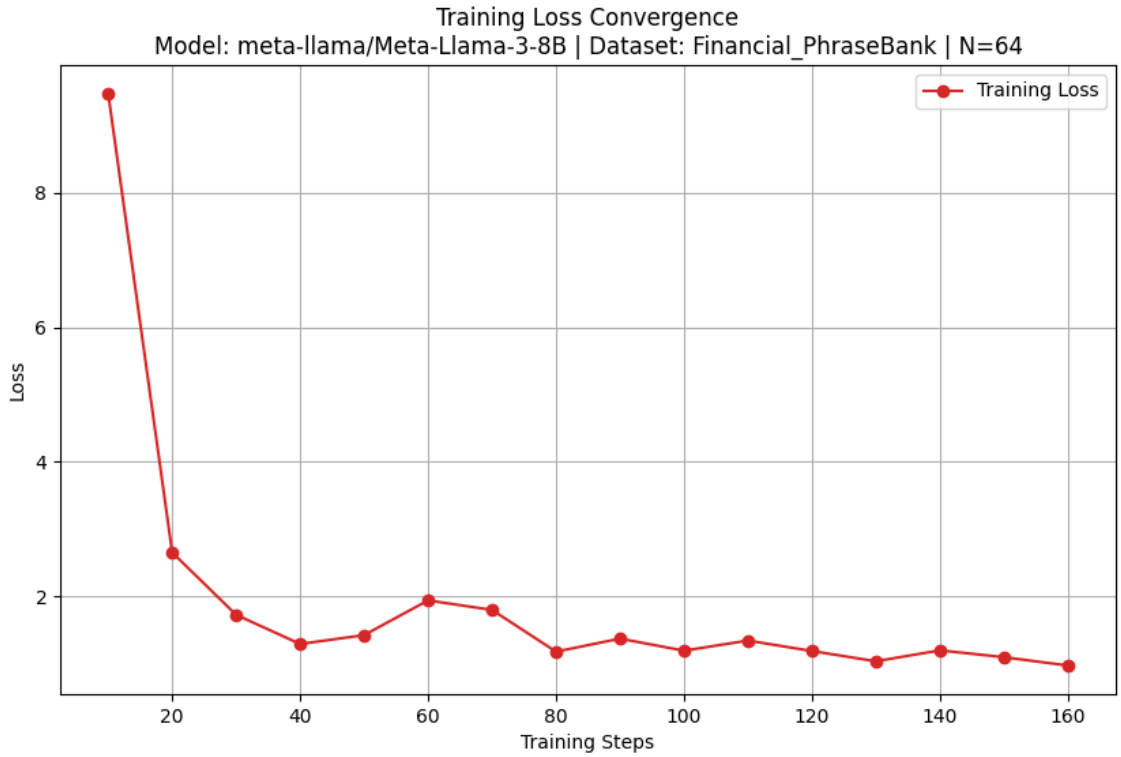
Training Loss Convergence
Model: meta-llama/Meta-Llama-3-8B | Dataset: Financial_PhraseBank | N=64
Training Loss
Loss
Training Steps
2
4
6
8
20
40
60
80
100
120
140
160

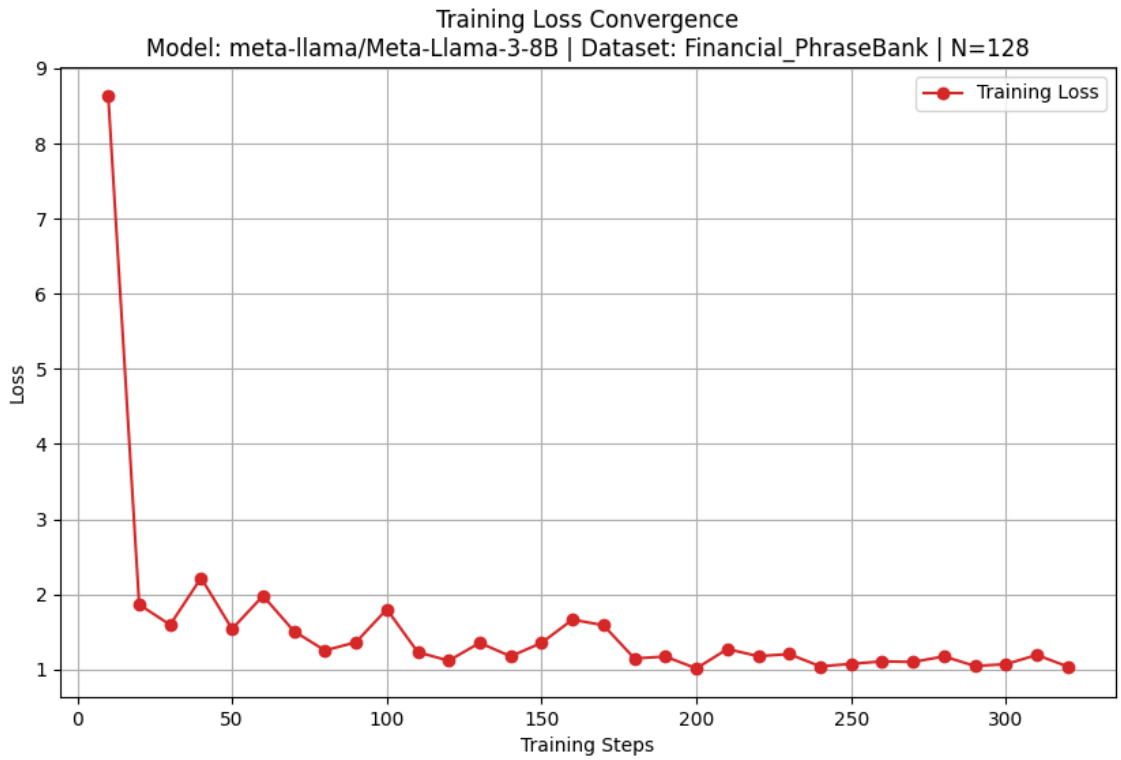
Training Loss Convergence
Model: meta-llama/Meta-Llama-3-8B | Dataset: Financial_PhraseBank | N=128
Training Loss
Loss
Training Steps
1
2
3
4
5
6
7
8
9
0
50
100
150
200
250
300

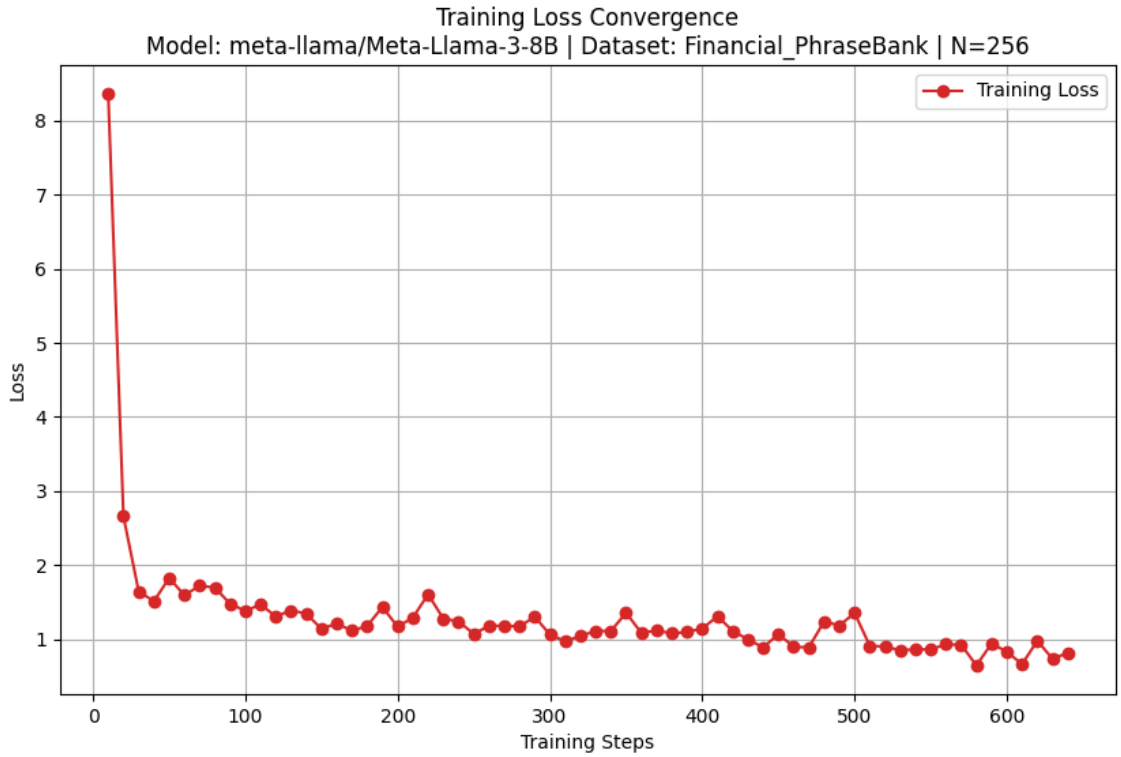
Training Loss Convergence
Model: meta-llama/Meta-Llama-3-8B | Dataset: Financial_PhraseBank | N=256
Training Loss
Loss
Training Steps
0
100
200
300
400
500
600
1
2
3
4
5
6
7
8

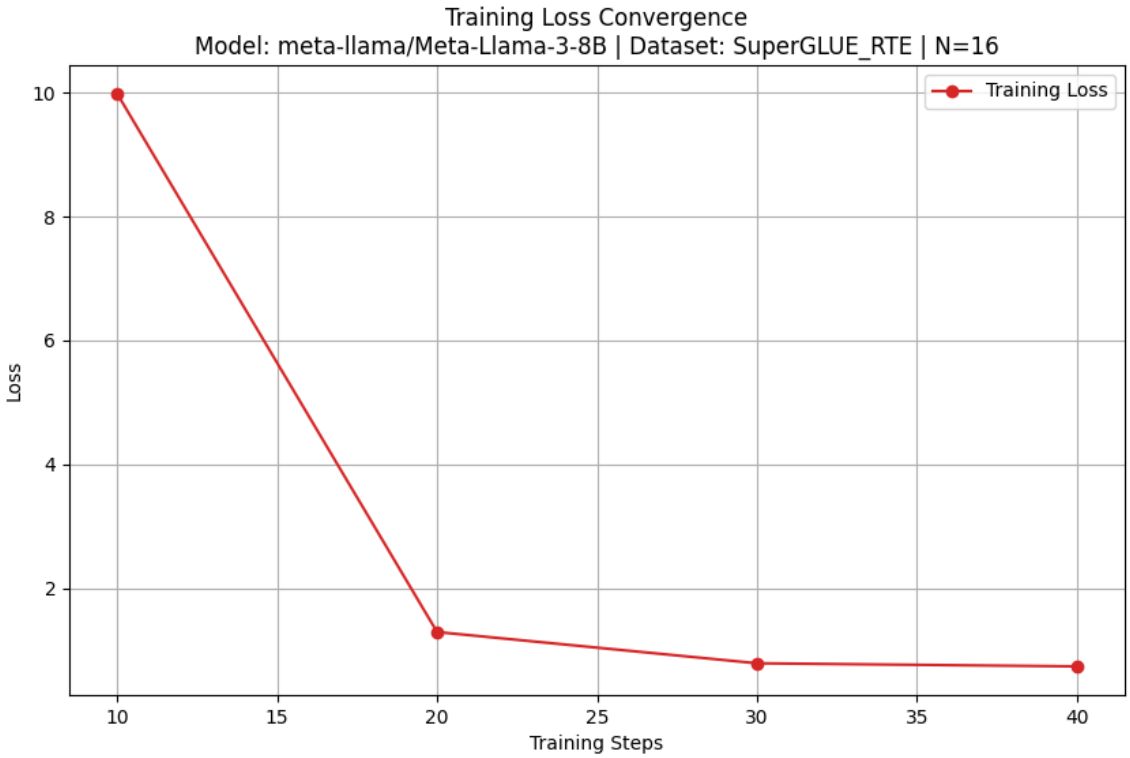
Training Loss Convergence
Model: meta-llama/Meta-Llama-3-8B | Dataset: SuperGLUE_RTE | N=16
Training Loss
Loss
Training Steps
10
15
20
25
30
35
40
2
4
6
8
10

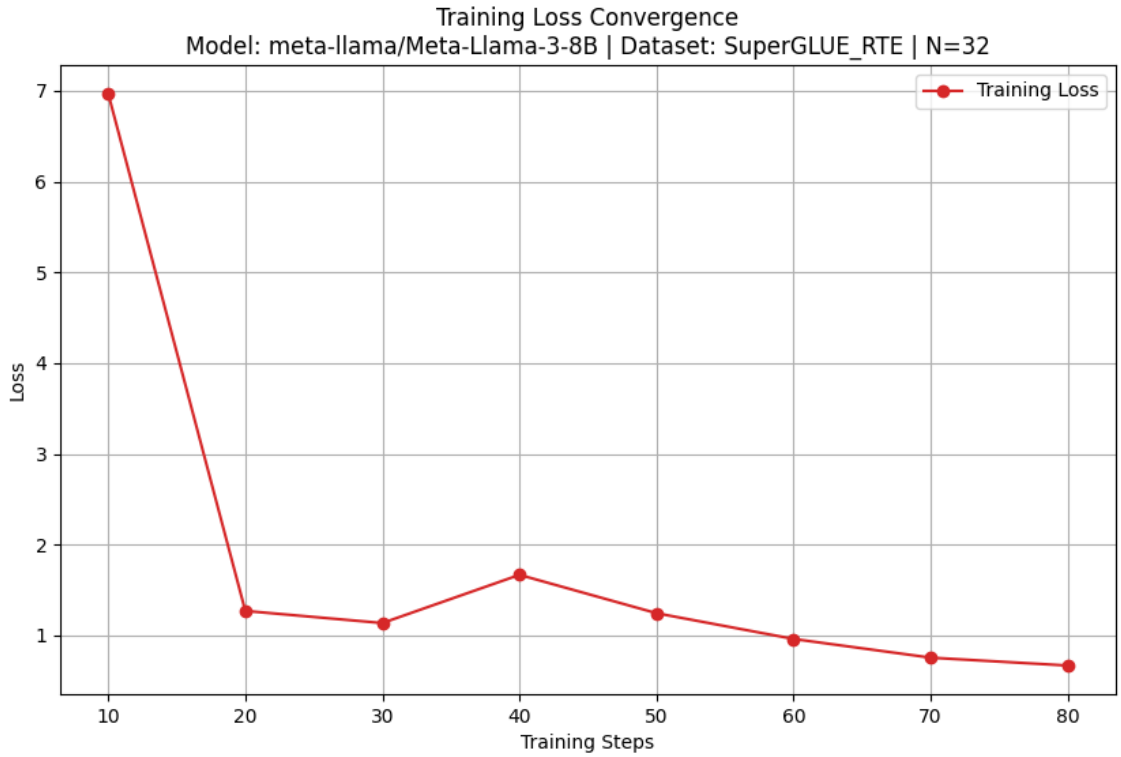
Training Loss Convergence
Model: meta-llama/Meta-Llama-3-8B | Dataset: SuperGLUE_RTE | N=32
Training Loss
Loss
Training Steps
7
6
5
4
3
2
1
10
20
30
40
50
60
70
80

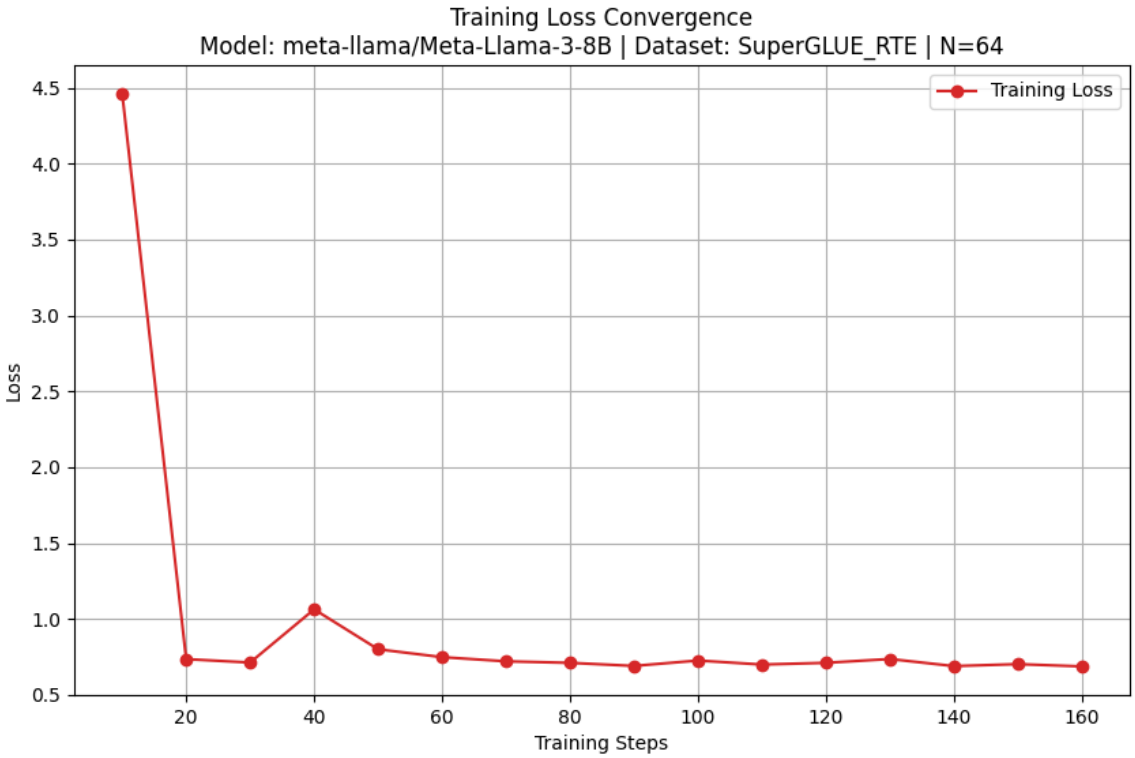
Training Loss Convergence
Model: meta-llama/Meta-Llama-3-8B | Dataset: SuperGLUE_RTE | N=64
Training Loss
Loss
Training Steps
4.5
4.0
3.5
3.0
2.5
2.0
1.5
1.0
0.5
20
40
60
80
100
120
140
160

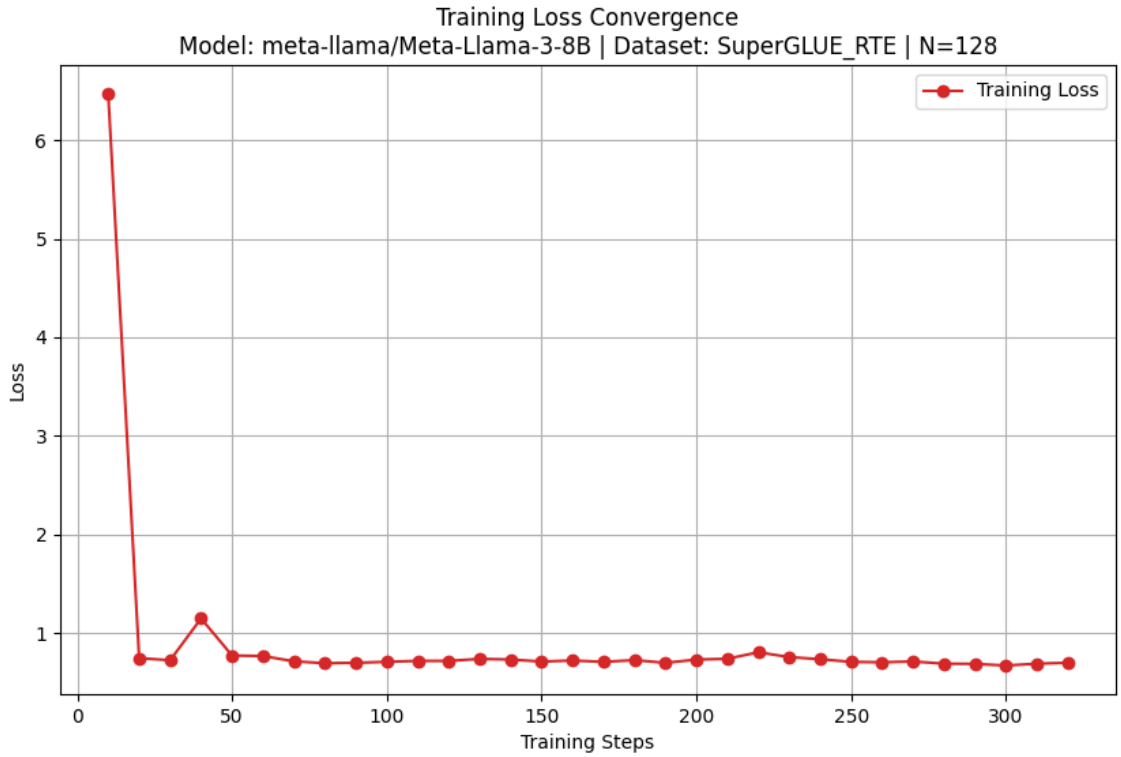
Training Loss Convergence
Model: meta-llama/Meta-Llama-3-8B | Dataset: SuperGLUE_RTE | N=128
Training Loss
Loss
Training Steps
0
50
100
150
200
250
300
1
2
3
4
5
6

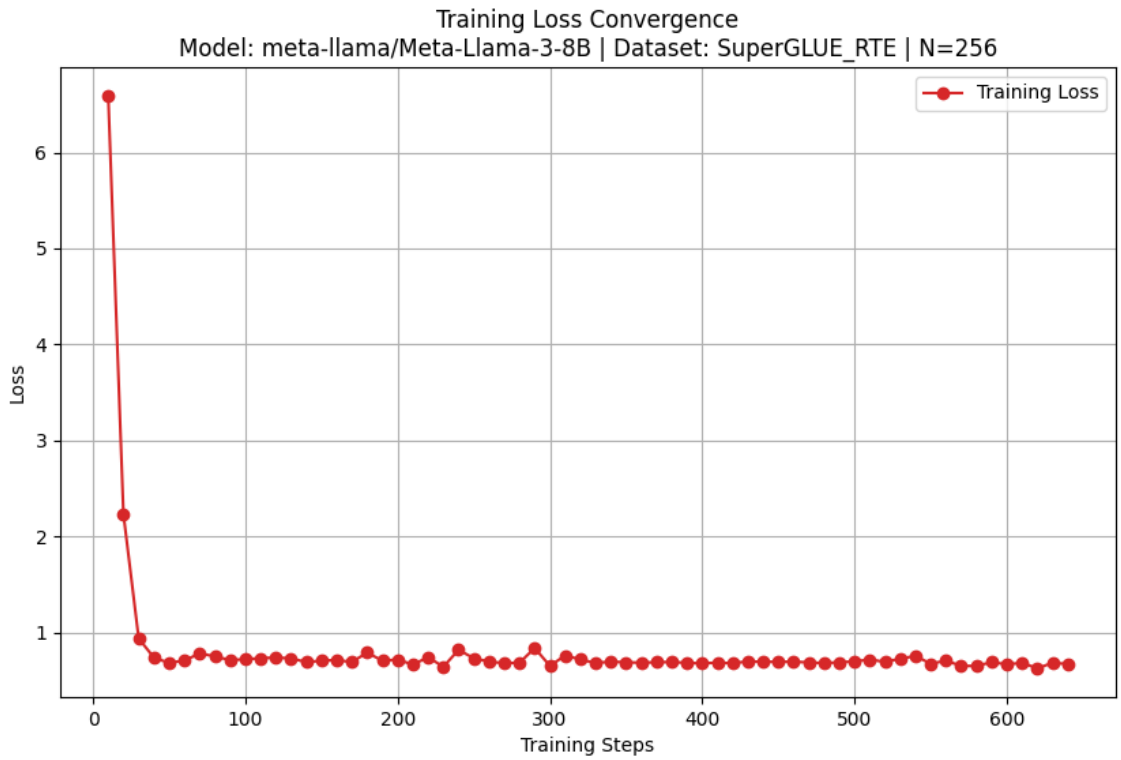
Training Loss Convergence
Model: meta-llama/Meta-Llama-3-8B | Dataset: SuperGLUE_RTE | N=256
Training Loss
Loss
Training Steps
0
100
200
300
400
500
600
1
2
3
4
5
6

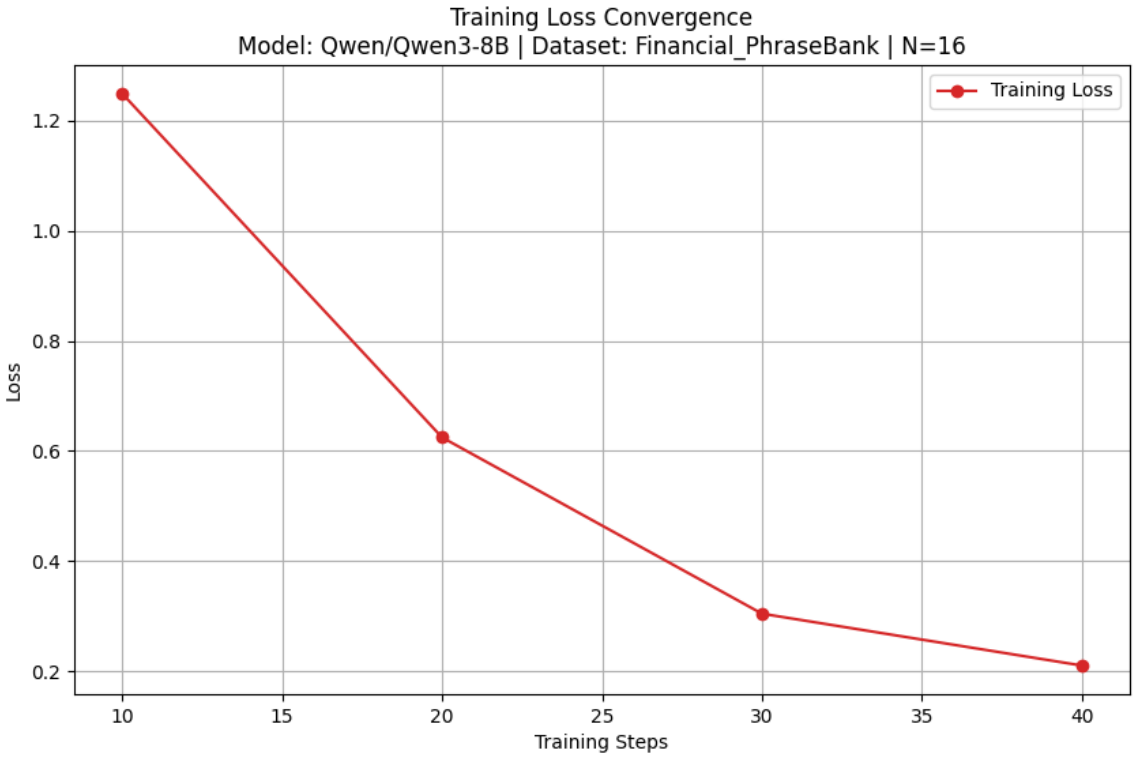
Training Loss Convergence
Model: Qwen/Qwen3-8B | Dataset: Financial_PhraseBank | N=16
Training Loss
Loss
Training Steps
1.2
1.0
0.8
0.6
0.4
0.2
10
15
20
25
30
35
40

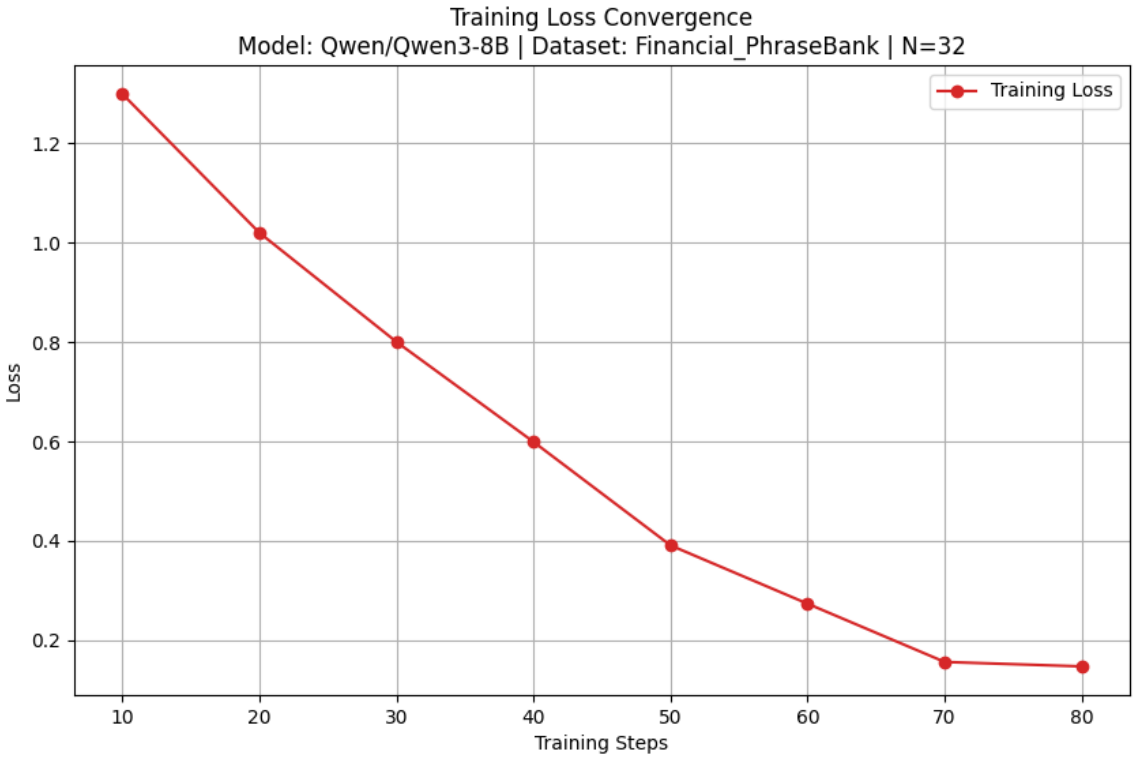
Training Loss Convergence
Model: Qwen/Qwen3-8B | Dataset: Financial_PhraseBank | N=32
Training Loss
Loss
Training Steps
1.2
1.0
0.8
0.6
0.4
0.2
10
20
30
40
50
60
70
80

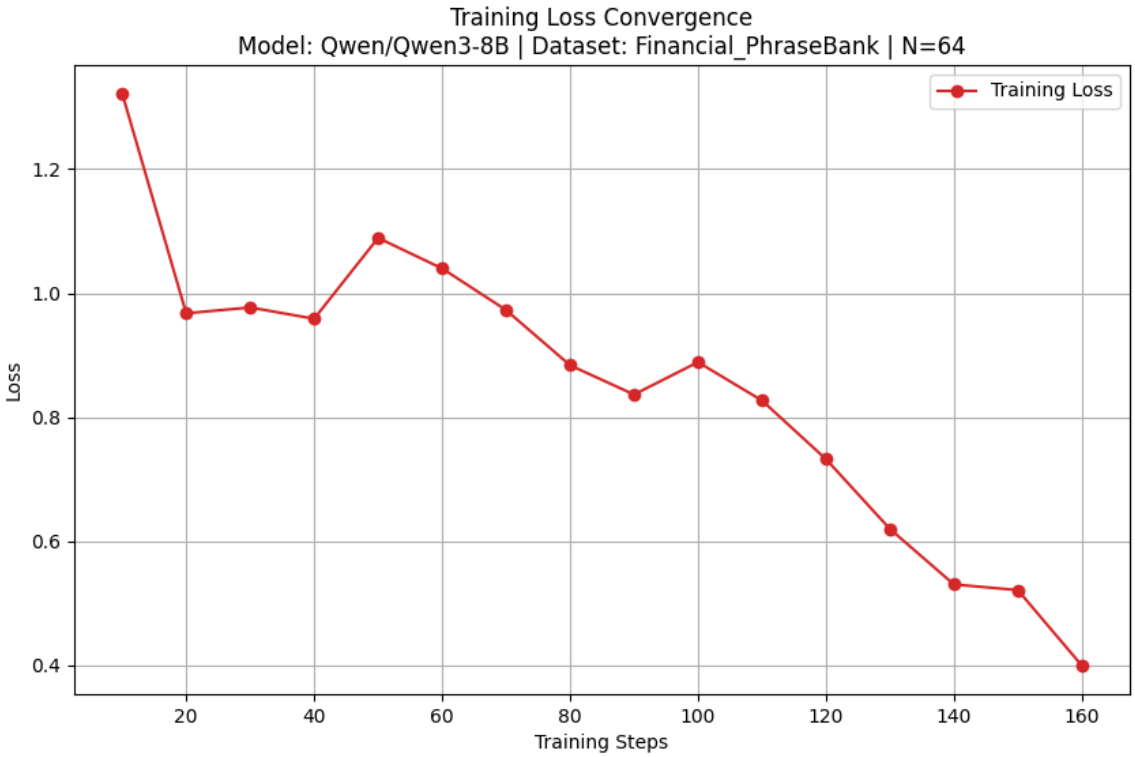
Training Loss Convergence
Model: Qwen/Qwen3-8B | Dataset: Financial_PhraseBank | N=64
Training Loss
Loss
Training Steps
1.2
1.0
0.8
0.6
0.4
20
40
60
80
100
120
140
160

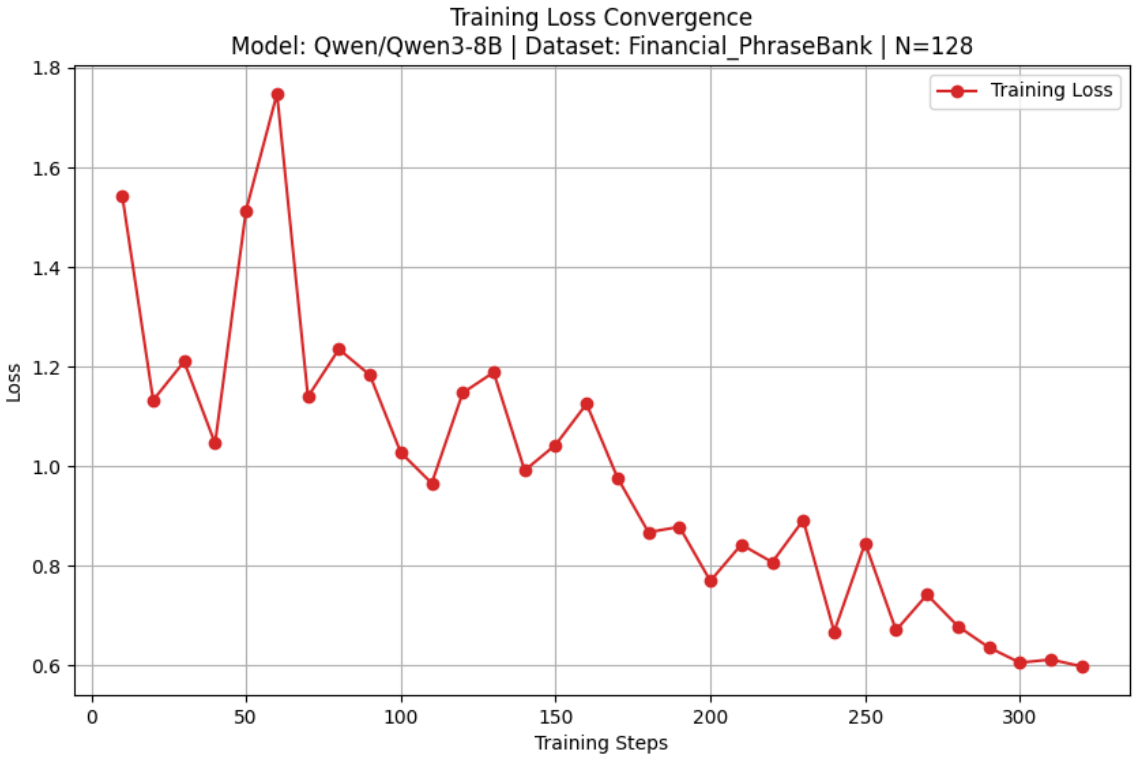
Training Loss Convergence
Model: Qwen/Qwen3-8B | Dataset: Financial_PhraseBank | N=128
Training Loss
Loss
Training Steps
1.8
1.6
1.4
1.2
1.0
0.8
0.6
0
50
100
150
200
250
300

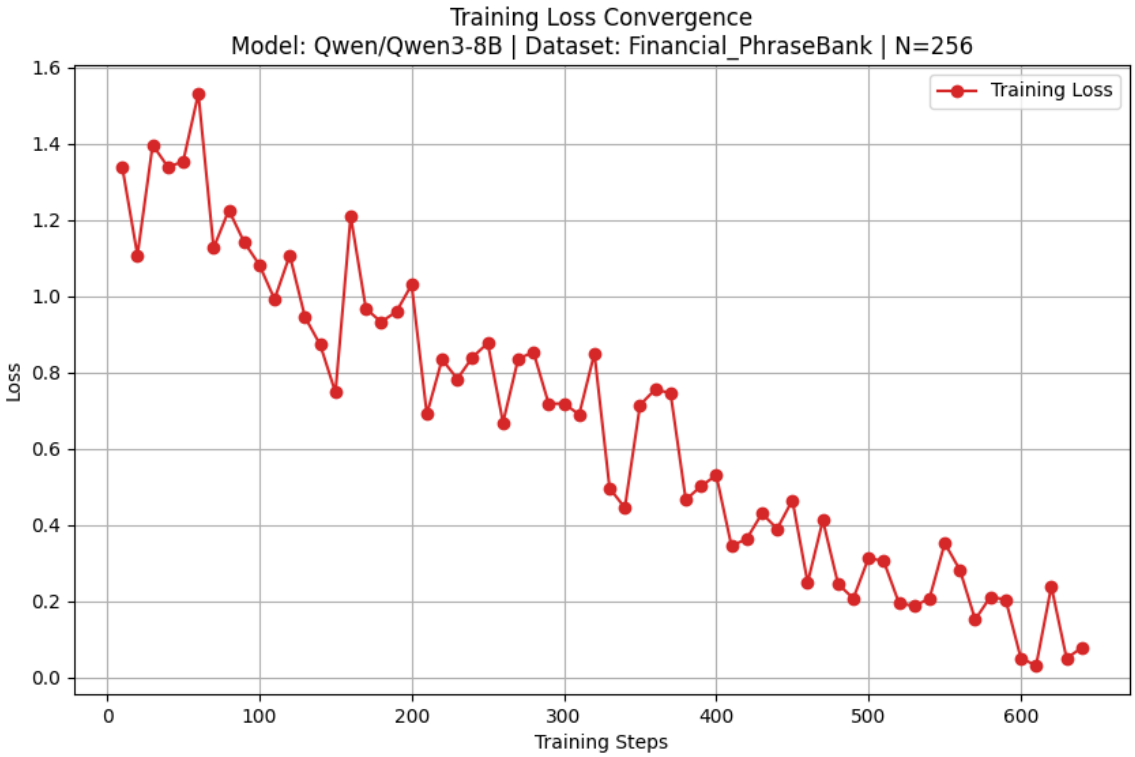
Training Loss Convergence
Model: Qwen/Qwen3-8B | Dataset: Financial_PhraseBank | N=256
Training Loss
Loss
Training Steps
0.0
0.2
0.4
0.6
0.8
1.0
1.2
1.4
1.6
0
100
200
300
400
500
600

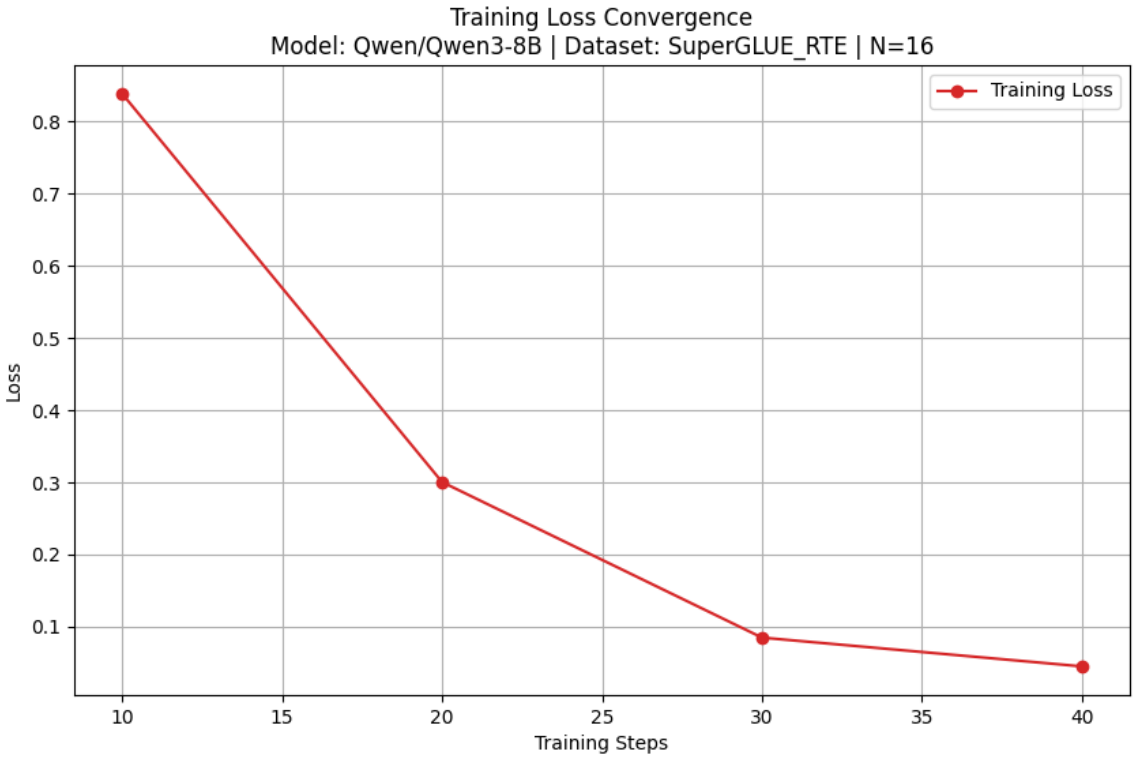
Training Loss Convergence
Model: Qwen/Qwen3-8B | Dataset: SuperGLUE_RTE | N=16
Training Loss
Loss
Training Steps
0.1
0.2
0.3
0.4
0.5
0.6
0.7
0.8
10
15
20
25
30
35
40

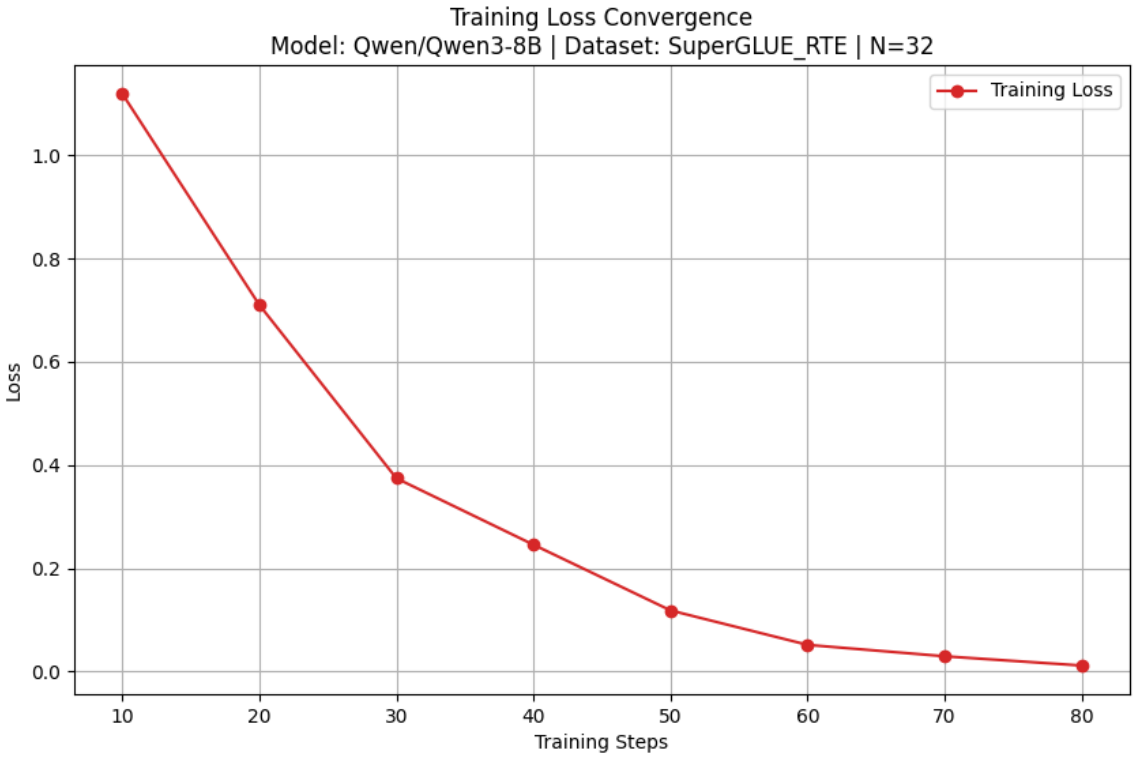
Training Loss Convergence
Model: Qwen/Qwen3-8B | Dataset: SuperGLUE_RTE | N=32
Training Loss
Loss
Training Steps
0.0
0.2
0.4
0.6
0.8
1.0
10
20
30
40
50
60
70
80

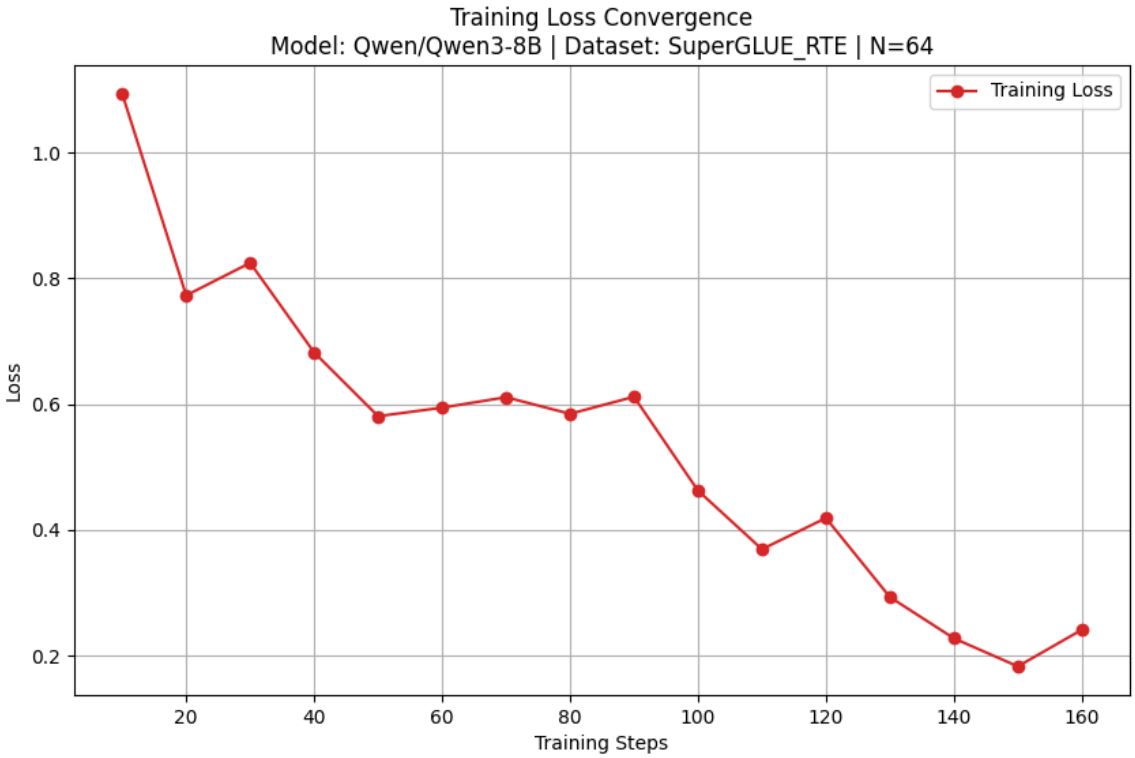
Training Loss Convergence
Model: Qwen/Qwen3-8B | Dataset: SuperGLUE_RTE | N=64
Training Loss
Loss
Training Steps
0.2
0.4
0.6
0.8
1.0
20
40
60
80
100
120
140
160

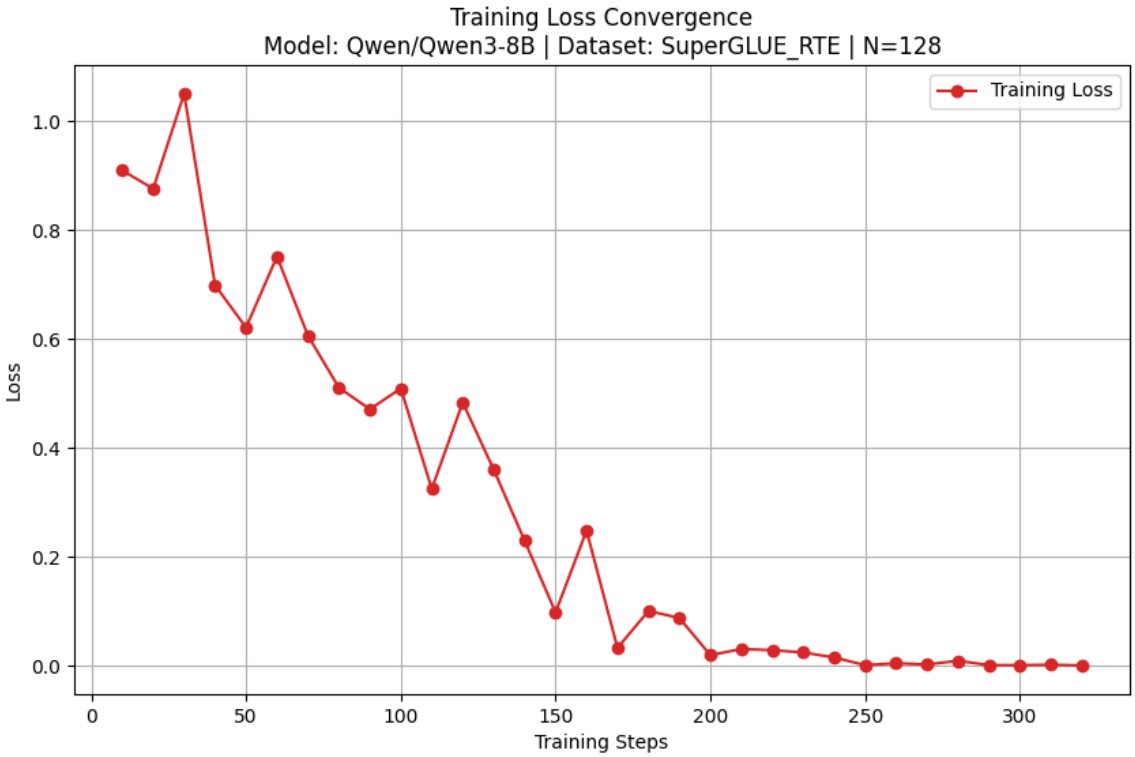
Training Loss Convergence
Model: Qwen/Qwen3-8B | Dataset: SuperGLUE_RTE | N=128
Training Loss
Loss
Training Steps
0.0
0.2
0.4
0.6
0.8
1.0
0
50
100
150
200
250
300

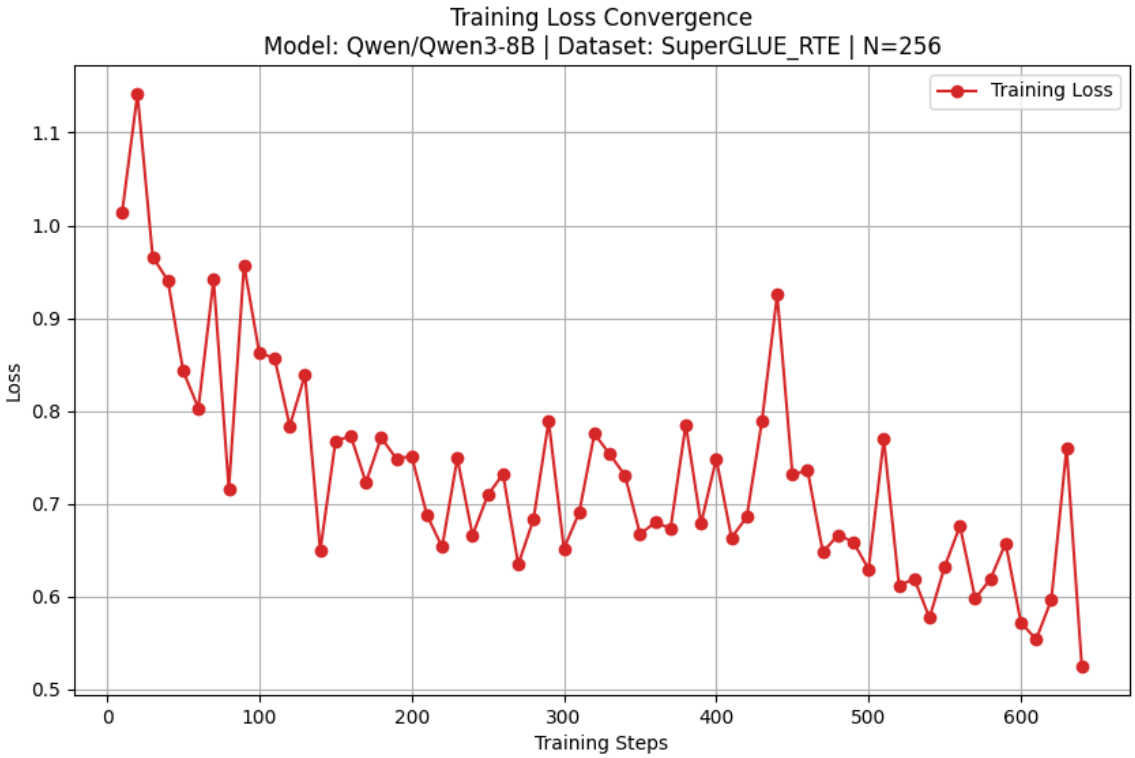
Training Loss Convergence
Model: Qwen/Qwen3-8B | Dataset: SuperGLUE_RTE | N=256
Training Loss
Loss
Training Steps
0.5
0.6
0.7
0.8
0.9
1.0
1.1
0
100
200
300
400
500
600