Step-by-Step Evaluation Metrics

Model: meta-llama/Meta-Llama-3-8B | Dataset: Financial_PhraseBank | N: 16

Acc: 0.1611 | Latency: 7.16ms | VRAM: 14.28GB

Model: meta-llama/Meta-Llama-3-8B | Dataset: Financial_PhraseBank | N: 32

Acc: 0.1611 | Latency: 7.15ms | VRAM: 14.28GB

Model: meta-llama/Meta-Llama-3-8B | Dataset: Financial_PhraseBank | N: 64

Acc: 0.1611 | Latency: 7.14ms | VRAM: 14.29GB

Model: meta-llama/Meta-Llama-3-8B | Dataset: Financial_PhraseBank | N: 128

Acc: 0.1611 | Latency: 7.16ms | VRAM: 14.29GB

Model: meta-llama/Meta-Llama-3-8B | Dataset: Financial_PhraseBank | N: 256

Acc: 0.1611 | Latency: 7.16ms | VRAM: 14.29GB

Model: meta-llama/Meta-Llama-3-8B | Dataset: SuperGLUE_RTE | N: 16

Acc: 0.5442 | Latency: 7.16ms | VRAM: 14.29GB

Model: meta-llama/Meta-Llama-3-8B | Dataset: SuperGLUE_RTE | N: 32

Acc: 0.5442 | Latency: 7.17ms | VRAM: 14.29GB

Model: meta-llama/Meta-Llama-3-8B | Dataset: SuperGLUE_RTE | N: 64

Acc: 0.5442 | Latency: 7.16ms | VRAM: 14.29GB

Model: meta-llama/Meta-Llama-3-8B | Dataset: SuperGLUE_RTE | N: 128

Acc: 0.5442 | Latency: 7.16ms | VRAM: 14.29GB

Model: meta-llama/Meta-Llama-3-8B | Dataset: SuperGLUE_RTE | N: 256

Acc: 0.5442 | Latency: 7.18ms | VRAM: 14.29GB

Model: Qwen/Qwen3-8B | Dataset: Financial_PhraseBank | N: 16

Acc: 0.1611 | Latency: 8.25ms | VRAM: 14.39GB

Model: Qwen/Qwen3-8B | Dataset: Financial_PhraseBank | N: 32

Acc: 0.1611 | Latency: 8.26ms | VRAM: 14.39GB

Model: Qwen/Qwen3-8B | Dataset: Financial_PhraseBank | N: 64

Acc: 0.1611 | Latency: 8.25ms | VRAM: 14.39GB

Model: Qwen/Qwen3-8B | Dataset: Financial_PhraseBank | N: 128

Acc: 0.1611 | Latency: 8.25ms | VRAM: 14.39GB

Model: Qwen/Qwen3-8B | Dataset: Financial_PhraseBank | N: 256

Acc: 0.1611 | Latency: 8.25ms | VRAM: 14.39GB

Model: Qwen/Qwen3-8B | Dataset: SuperGLUE_RTE | N: 16

Acc: 0.5442 | Latency: 8.26ms | VRAM: 14.39GB

Model: Qwen/Qwen3-8B | Dataset: SuperGLUE_RTE | N: 32

Acc: 0.5442 | Latency: 8.27ms | VRAM: 14.39GB

Model: Qwen/Qwen3-8B | Dataset: SuperGLUE_RTE | N: 64

Acc: 0.5442 | Latency: 8.27ms | VRAM: 14.39GB

Model: Qwen/Qwen3-8B | Dataset: SuperGLUE_RTE | N: 128

Acc: 0.5442 | Latency: 8.25ms | VRAM: 14.39GB

Model: Qwen/Qwen3-8B | Dataset: SuperGLUE_RTE | N: 256

Acc: 0.5442 | Latency: 8.25ms | VRAM: 14.39GB

Tradeoff Analysis: Financial_PhraseBank - Accuracy

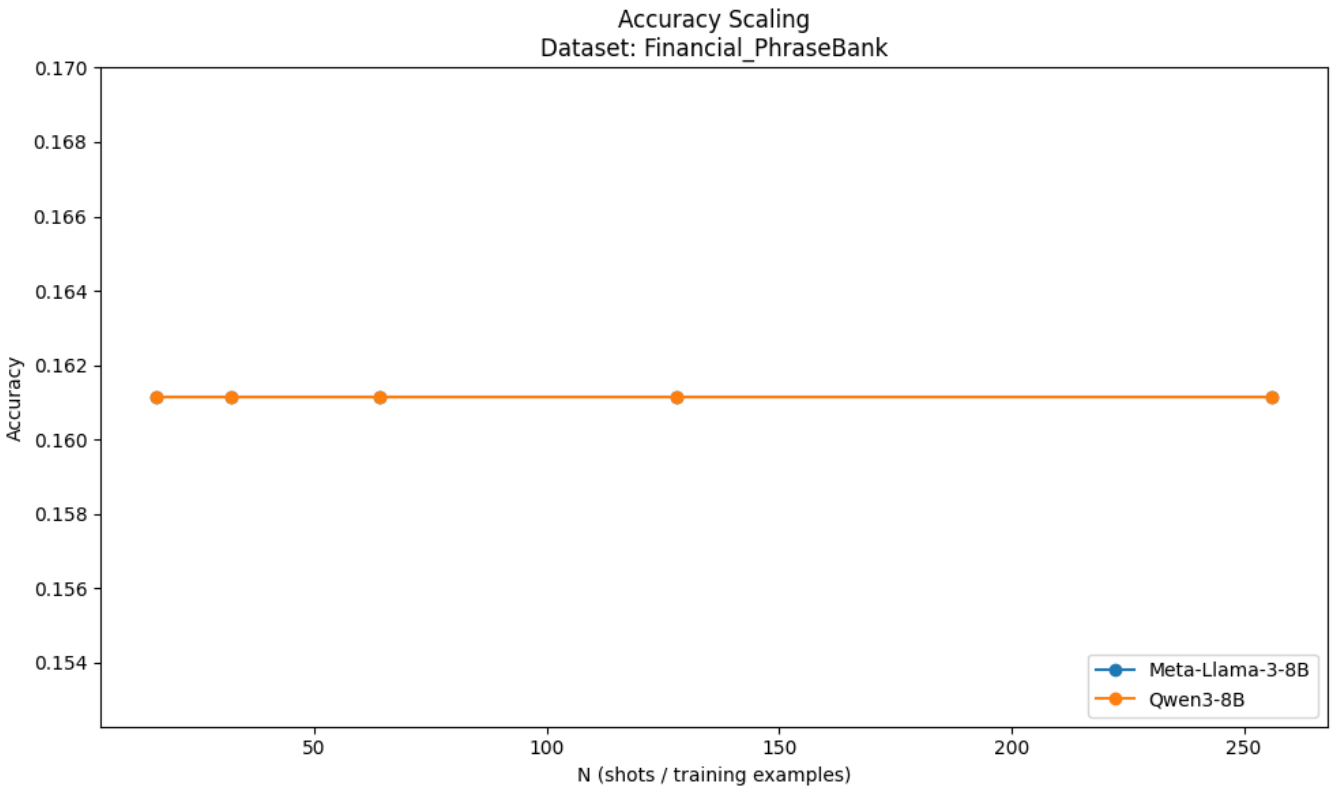

Tradeoff Analysis: Financial_PhraseBank - Latency (ms)

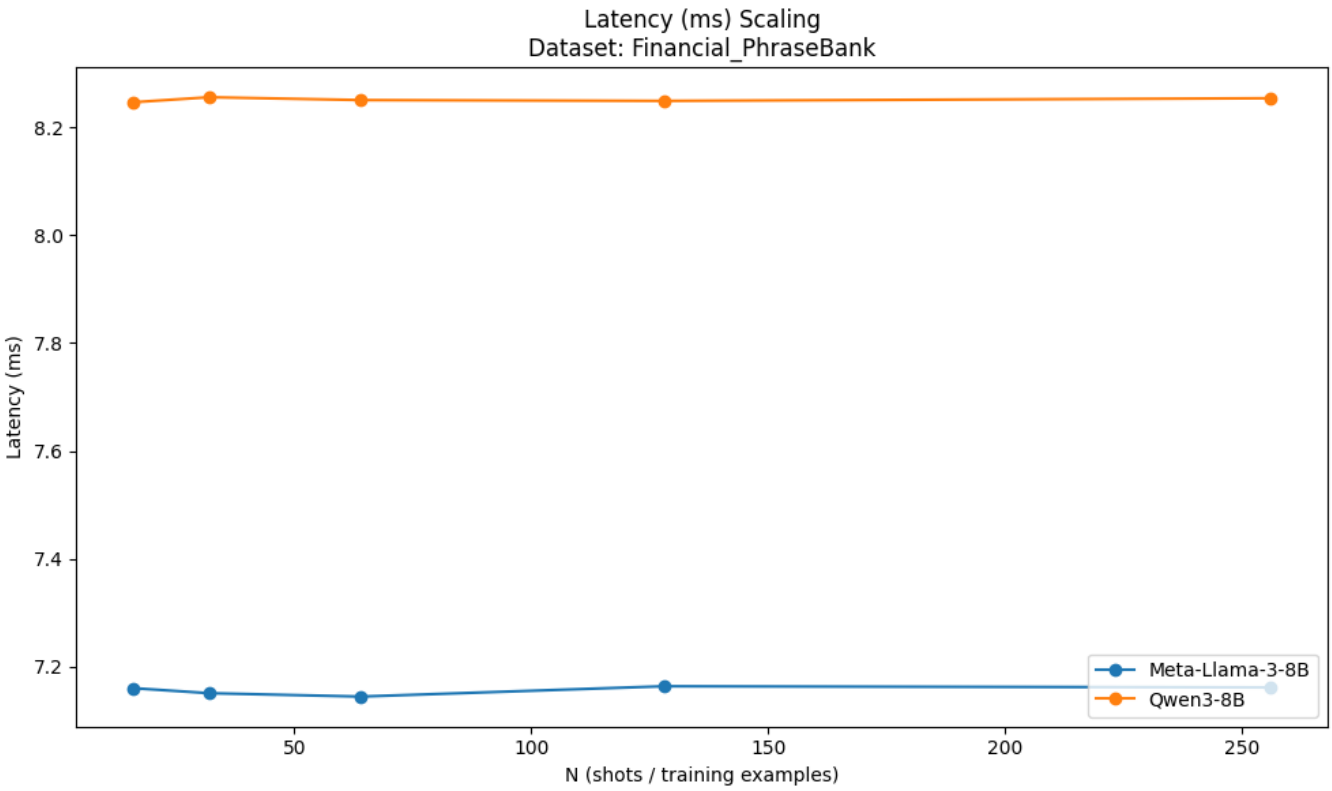

Tradeoff Analysis: Financial_PhraseBank - Throughput (samples/sec)

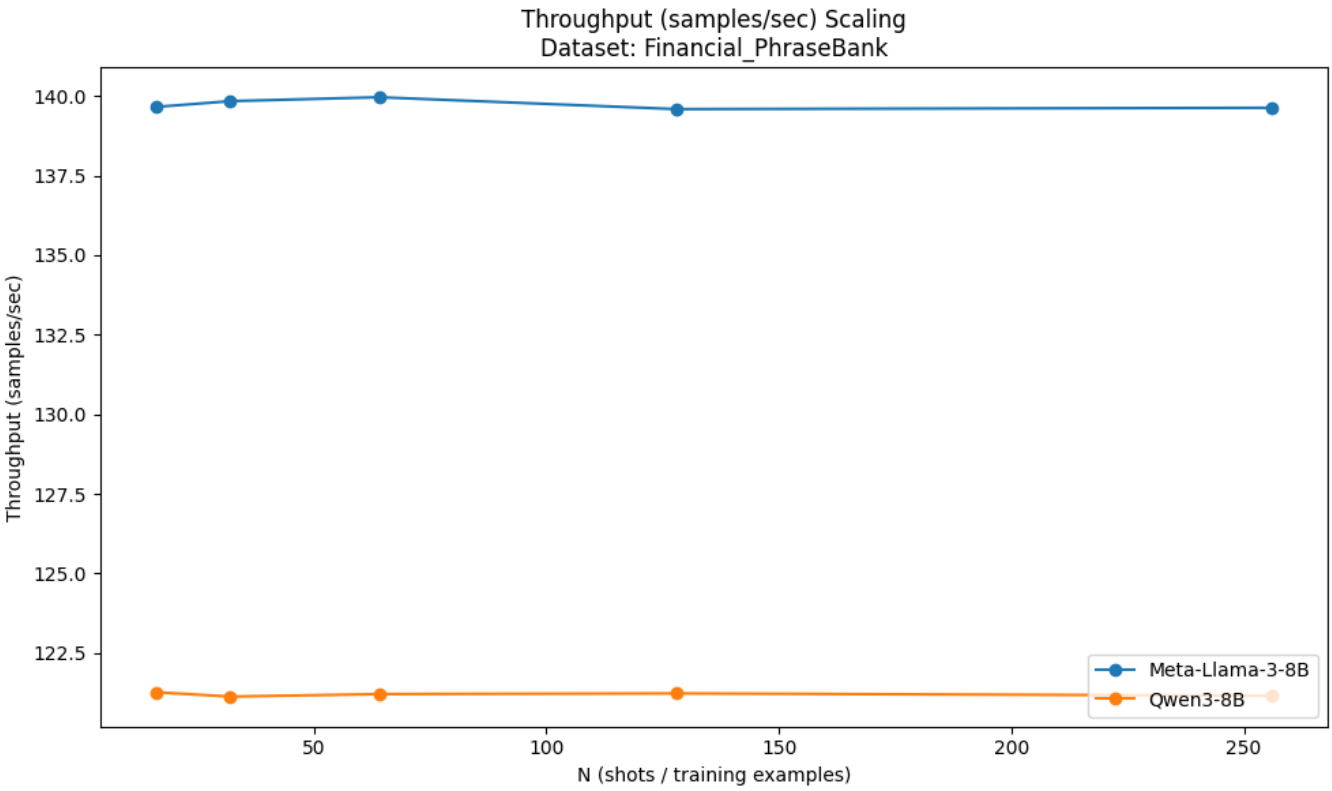

Tradeoff Analysis: Financial_PhraseBank - Peak VRAM (GB)

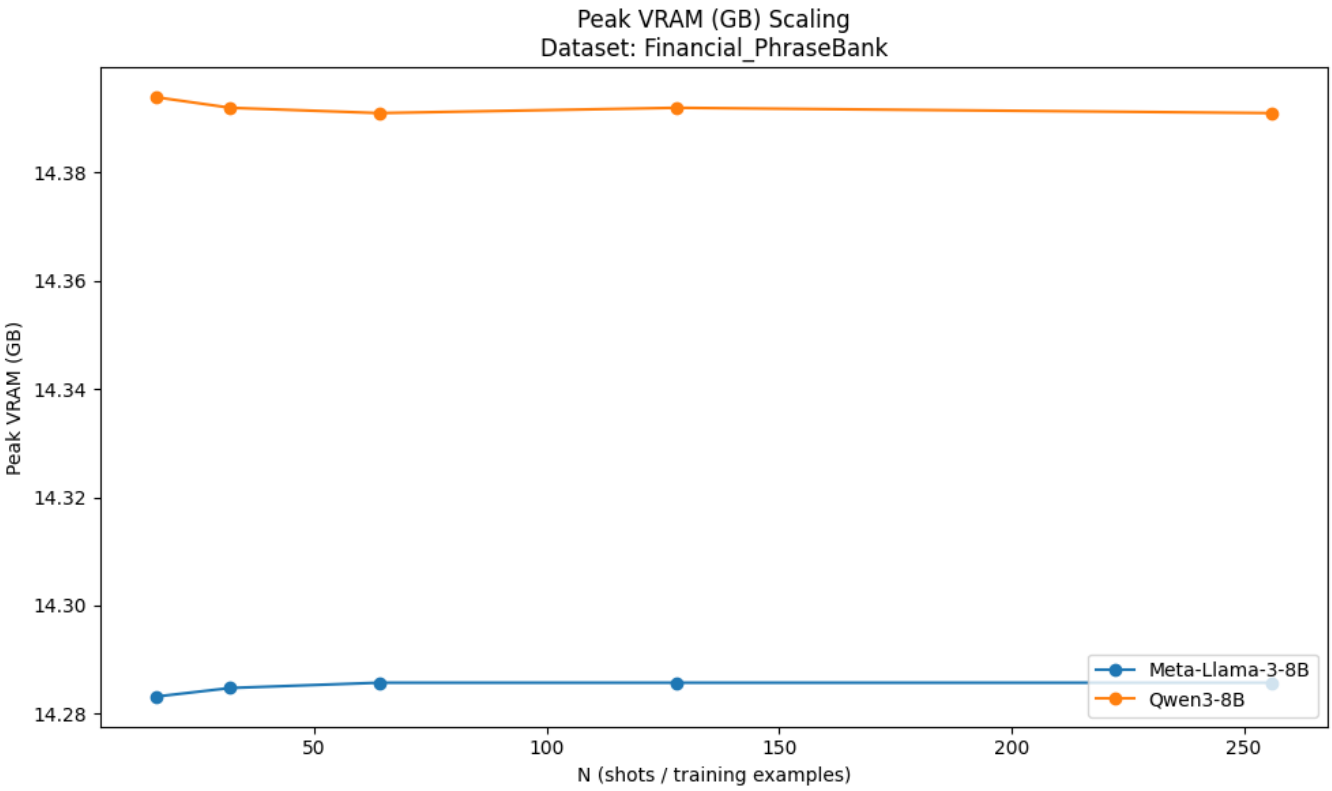

Tradeoff Analysis: SuperGLUE_RTE - Accuracy

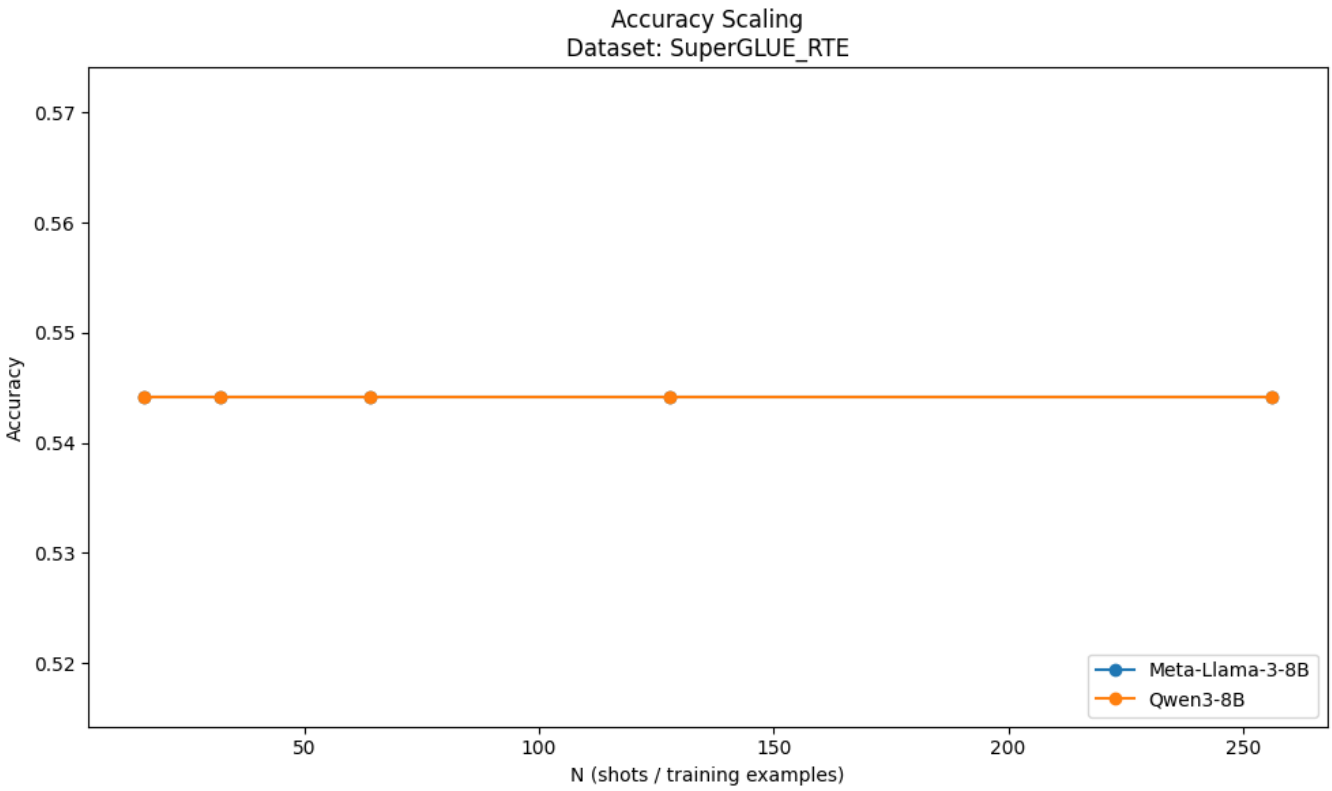

Tradeoff Analysis: SuperGLUE_RTE - Latency (ms)

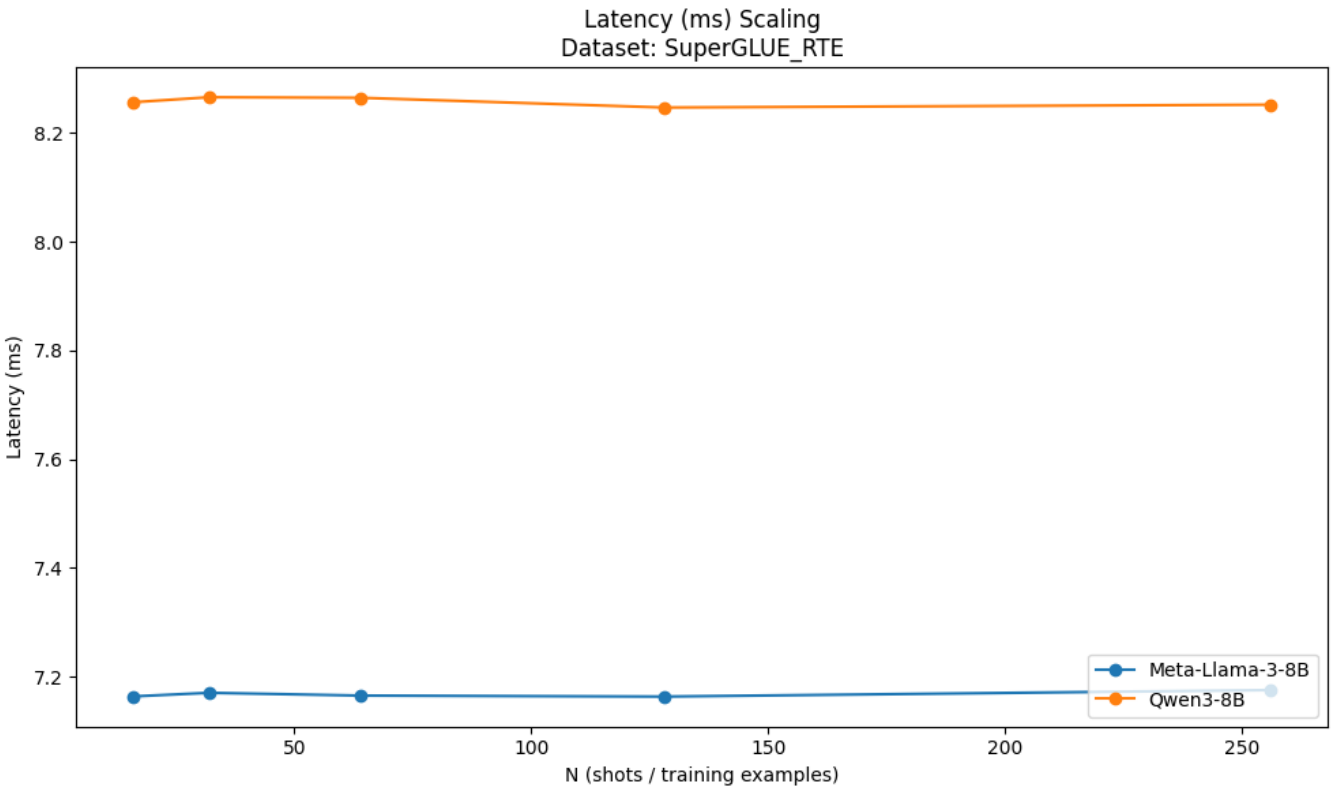

Tradeoff Analysis: SuperGLUE_RTE - Throughput (samples/sec)

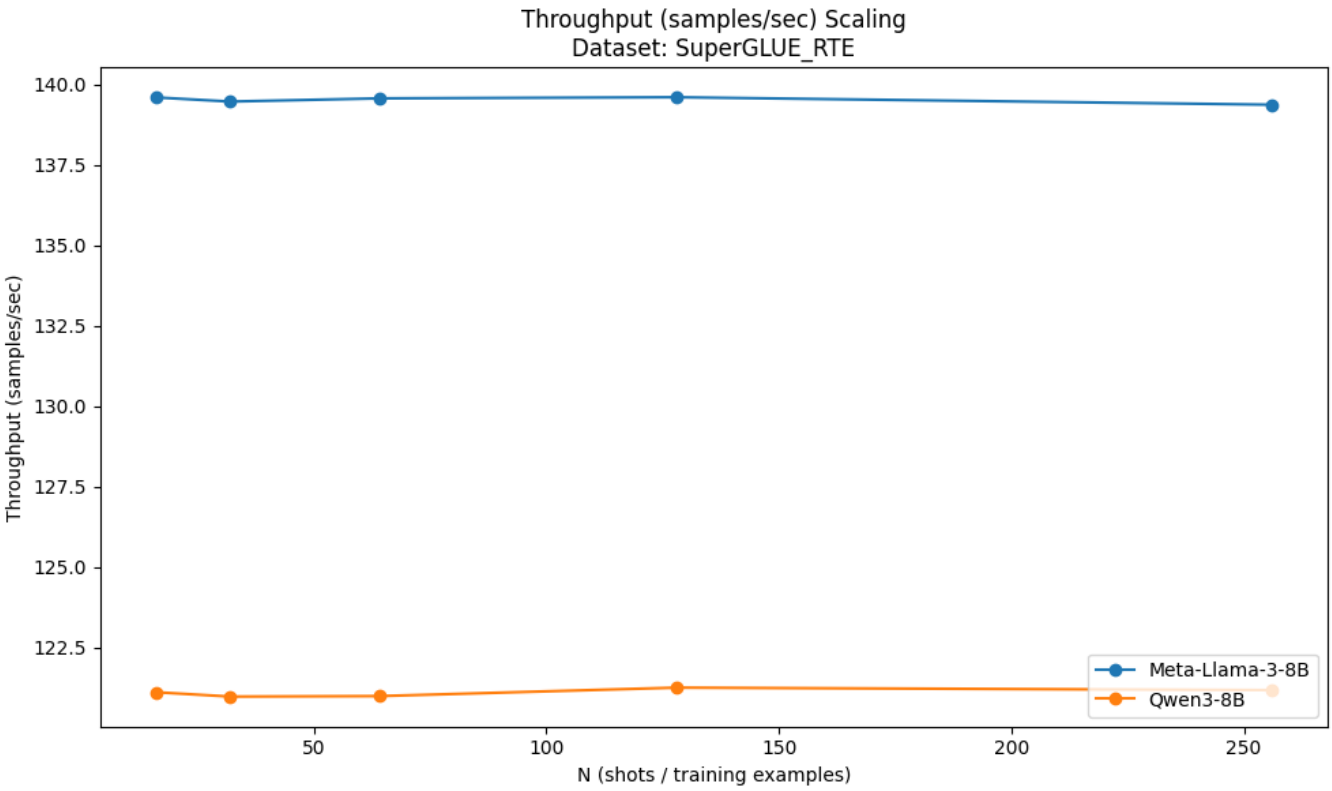
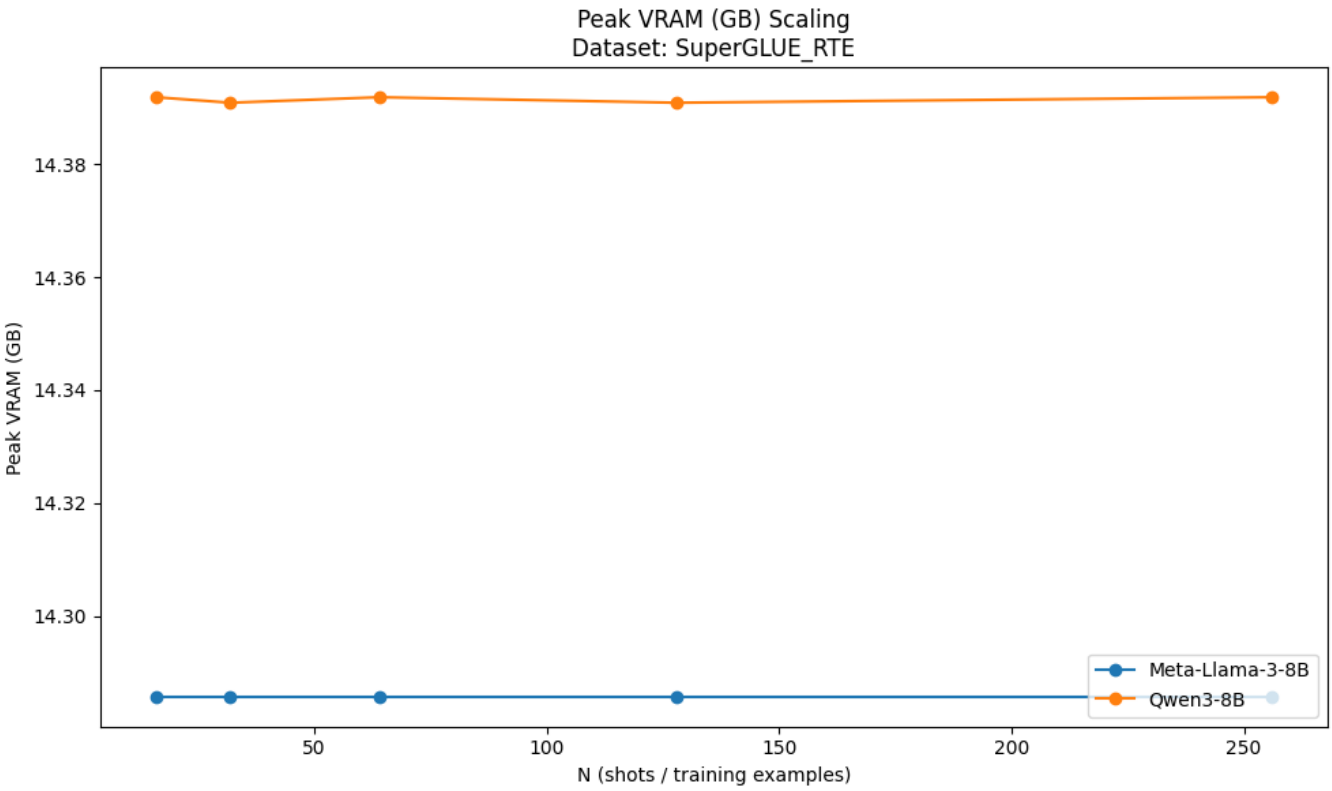
Tradeoff Analysis: SuperGLUE_RTE - Peak VRAM (GB)